



EE 044252: Digital Systems and Computer Structure
Spring 2019

Lecture 11: RISC-V Multi-Cycle



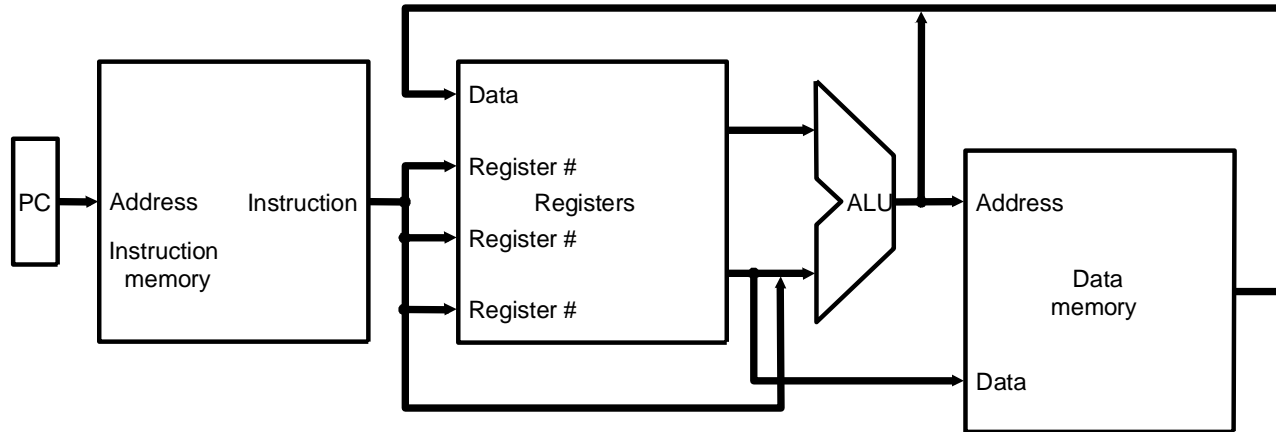
EE 044252: Digital Systems and Computer Structure

Topic	wk	Lectures	Tutorials	Workshop	Simulation
Arch	1	Intro. RISC-V architecture	Numbers. Codes		
Comb	2	Switching algebra & functions	Assembly programming		
	3	Combinational logic	Logic minimization	Combinational	
	4	Arithmetic. Memory	Gates		Combinational
Seq	5	Finite state machines	Logic		
	6	Sync FSM	Flip flops, FSM timing	Sequential	Sequential
	7	FSM equiv, scan, pipeline	FSM synthesis		
	8	Serial comm, RISC-V functions	Serial comm, pipeline		
μ Arch	9	RISC-V instruction formats	Function call		
	10	RISC-V single cycle	Single cycle RISC-V		Multi-cycle
	11	Multi-cycle RISC-V	Multi-cycle RISC-V		
	12	Interrupts, pipeline RISC-V	Microcode, interrupts		
	13	Dependencies in pipeline RISC-V	Depend. in pipeline RISC-V		

Agenda

- Multi-Cycle RISC-V Datapath
- Performance
- ROM Controller
- Micro-Coded Controller

Reminder: Single Cycle Datapath



More about single cycle datapath

- Simple to design and understand
- Performance limited by critical path—slowest instruction
 - LW
- Discouraged from adding even more complicated instructions

Elastic Cycle Time

- We saw that Load and Store instructions take more time than others
 - May be we can write a program that contains only “fast” commands?
- We know the dynamic instruction distribution (in many programs):
 - 24% - Loads, 12% - Stores, 44% - R-Type, 18% - Branches, 2% - Jumps
- We know the actual time required for each instruction (from Lecture 10):
 - Load—800ps, Stores—700ps, R-Type—600ps, Branches/Jumps—500ps
- Thus, the average time required for instruction execution is:
 - $800 \times 24\% + 700 \times 12\% + 600 \times 44\% + 500 \times 18\% + 500 \times 2\% = 640 \text{ ps}$
- We could create an elastic clock. It will allocate just the required time for each instruction
 - Its average cycle time is 640ps
 - Its average frequency is $1/640\text{ps} = 1.56\text{GHz}$
- Improvement relative to single cycle: $800\text{ps} / 640\text{ps} = 1.56\text{G} / 1.25\text{G} = 1.25$

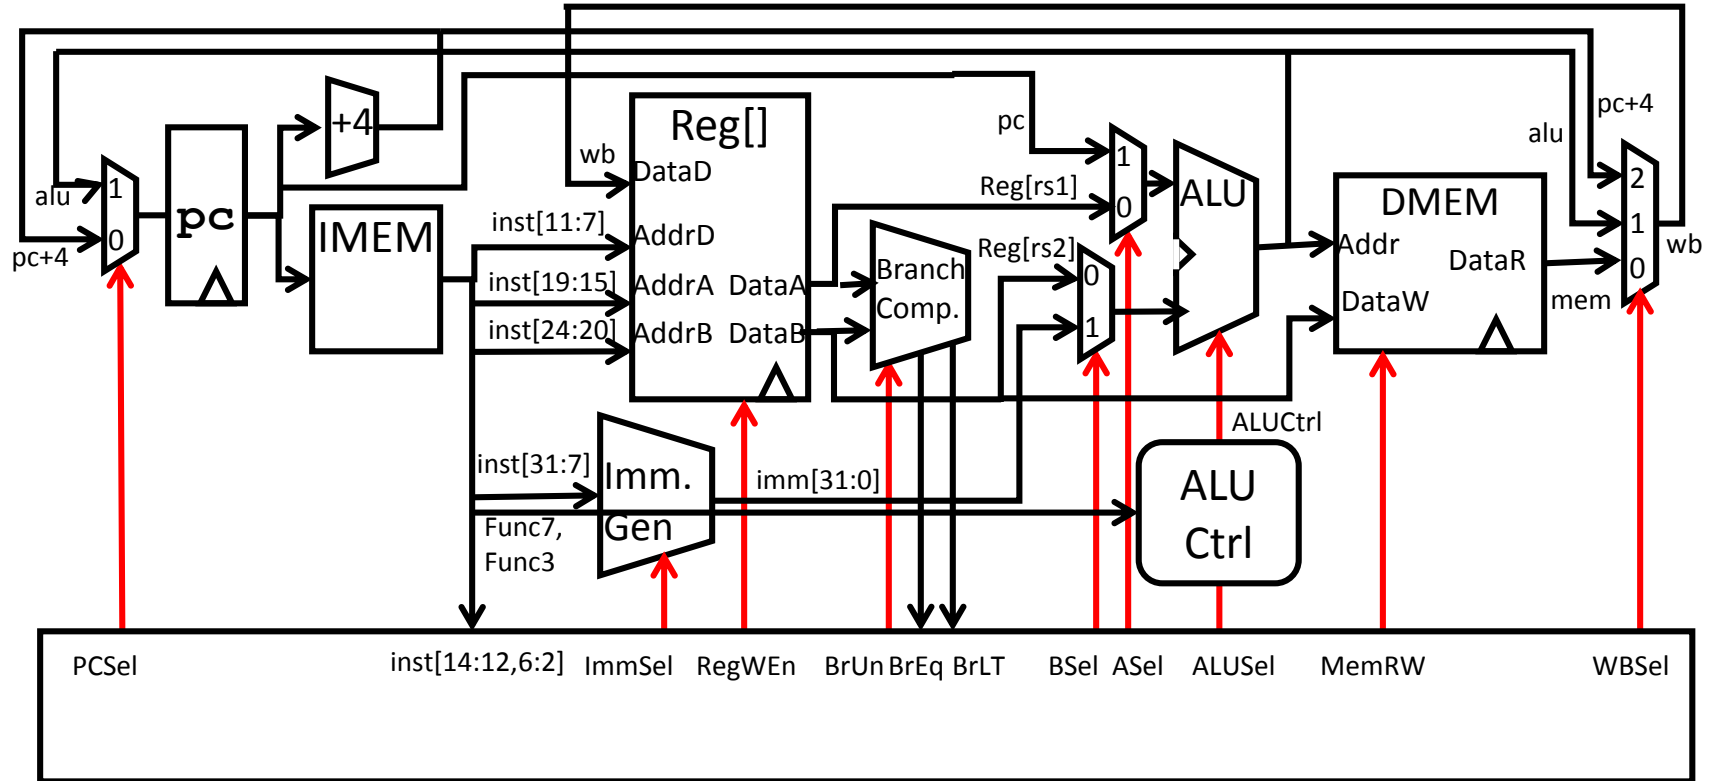
Multi-Cycle RISC-V

- Let's break computation down into multiple stages
- Perform one stage per one clock cycle
- Different instructions require a different number of cycles
- Goal: Faster computer. How is this possible?
- Same functional units can be re-used in multiple cycles of same instruction
 - We could not do that in a single-cycle RISC-V
- Need new registers for intermediate storage (in-between cycles of same instruction)
- Single-cycle RISC-V uses a combinational controller
Multi-cycle RISC-V uses a FSM controller

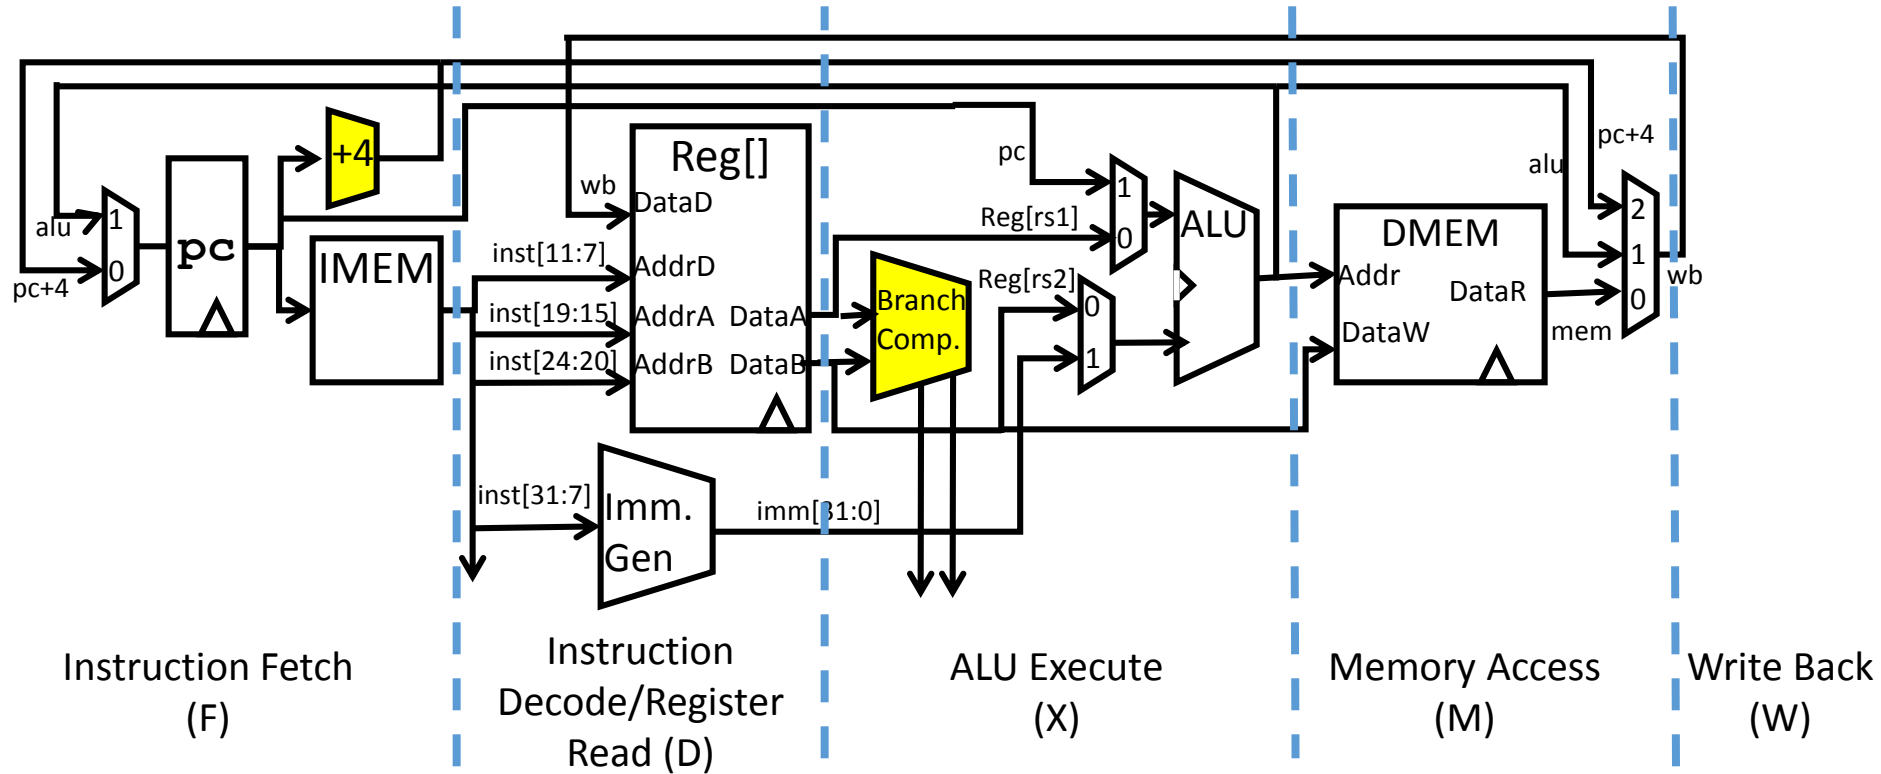
Two types of registers

- Registers holding persistent data
 - Contents used/written in one instruction can be used in a following instruction
 - Register File (32 registers), PC
 - Defined in architecture, visible in instruction simulators (e.g., venus)
- **NEW in multi-cycle:** Registers holding temporary data
 - Contents held only during execution of a single instruction
 - Unknown to programmer and to assembly language

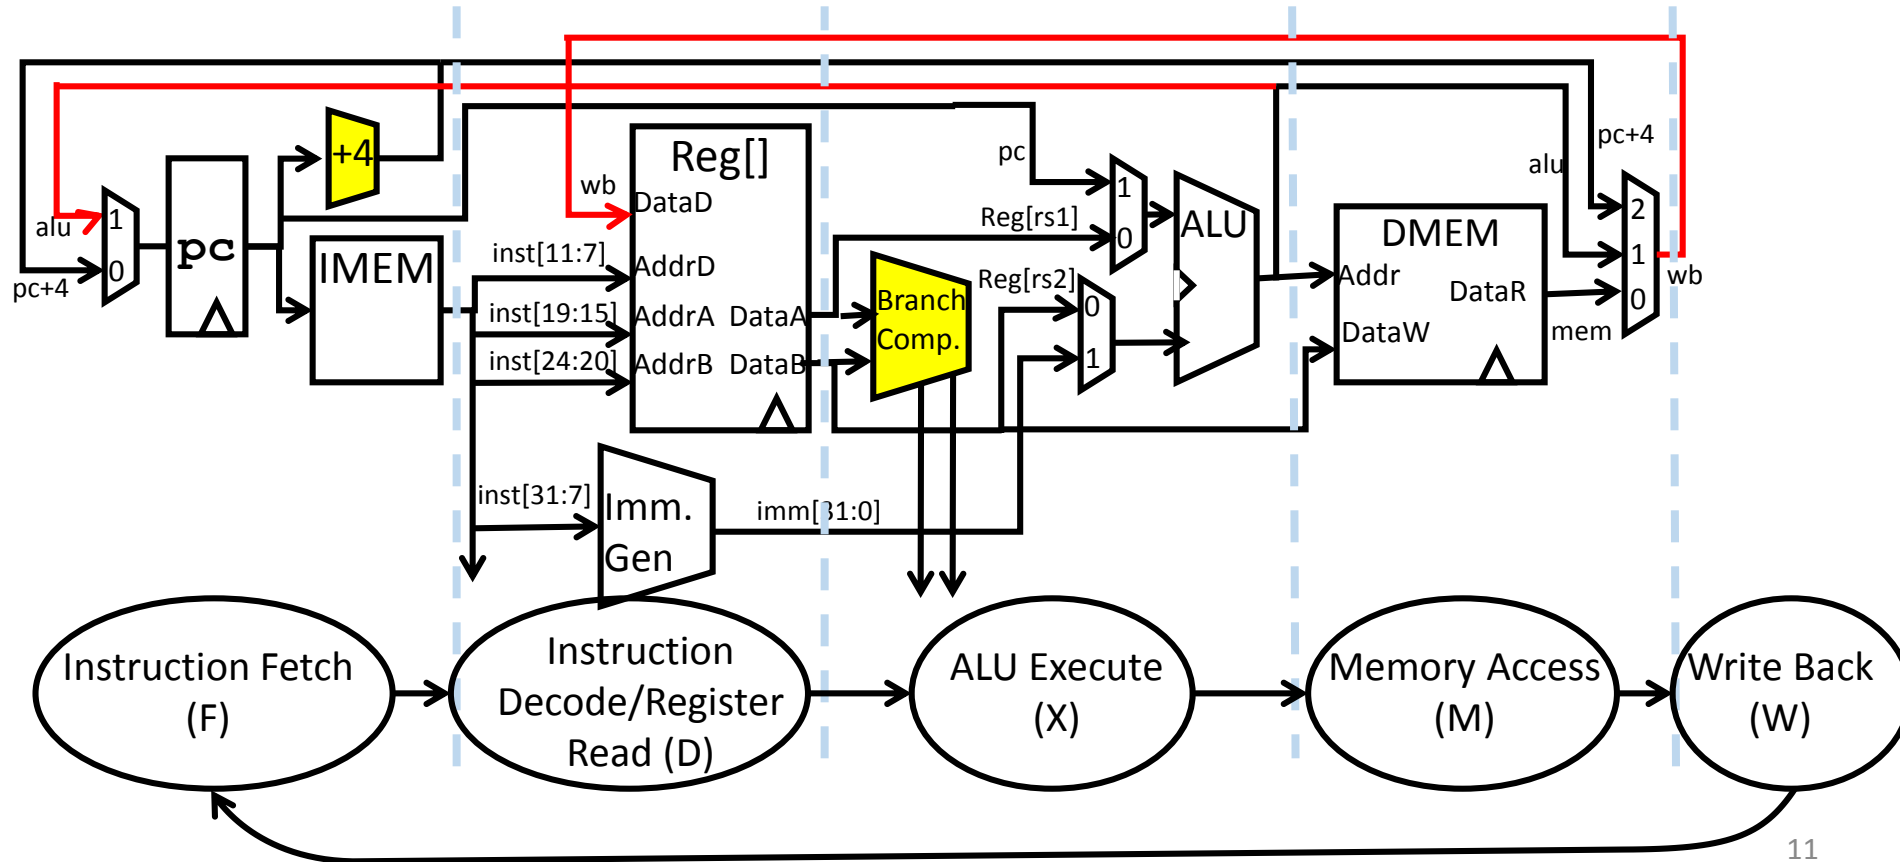
Single-Cycle RISC-V RV32I Datapath



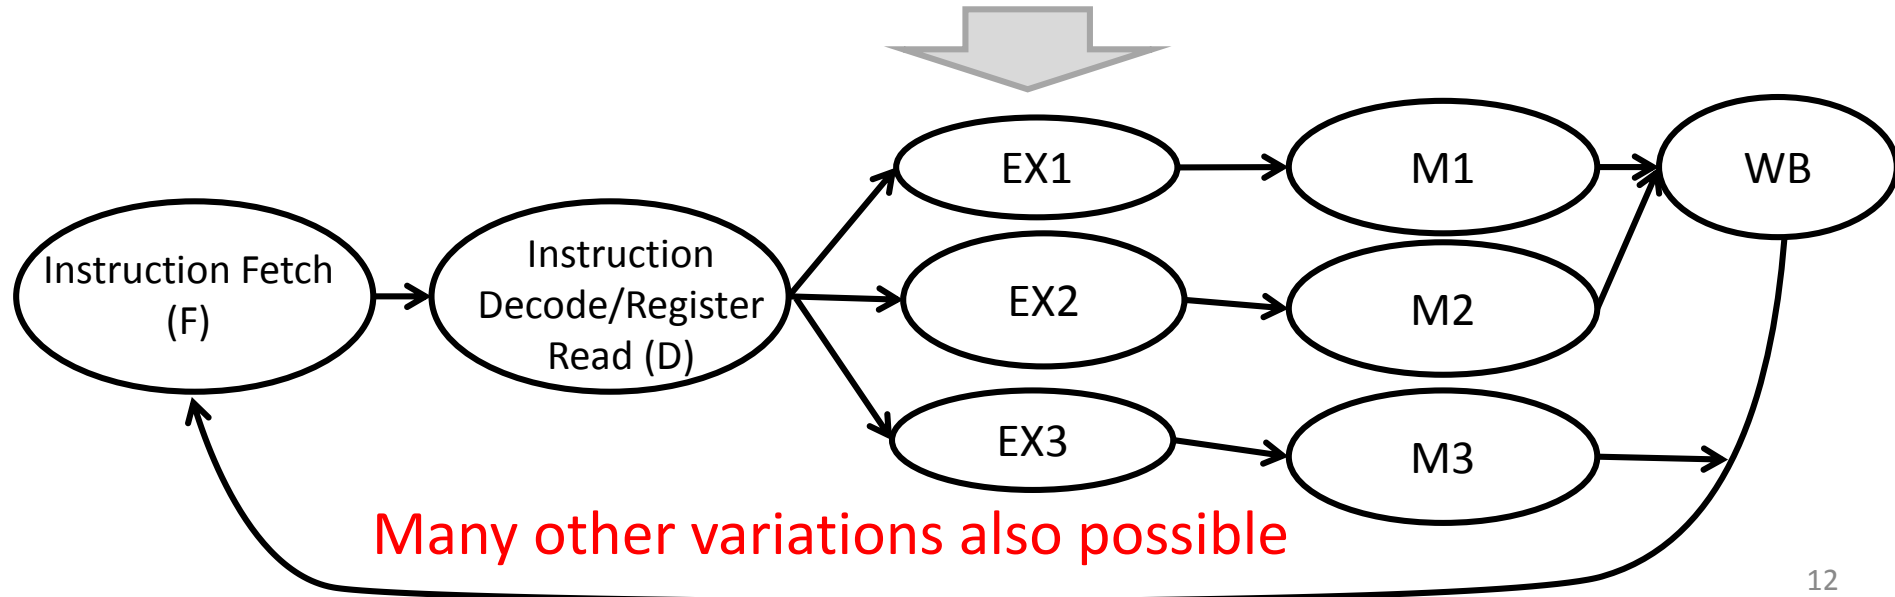
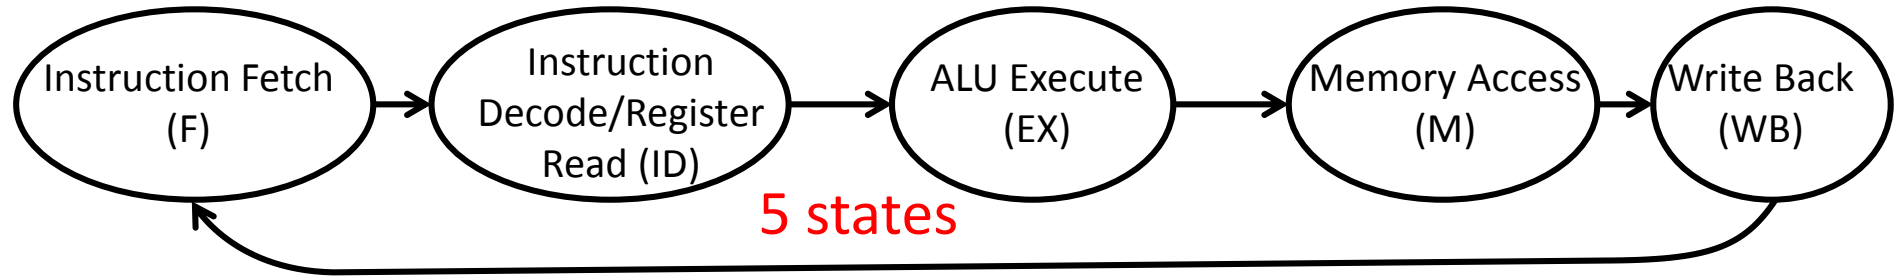
Breaking the datapath into 5 stages



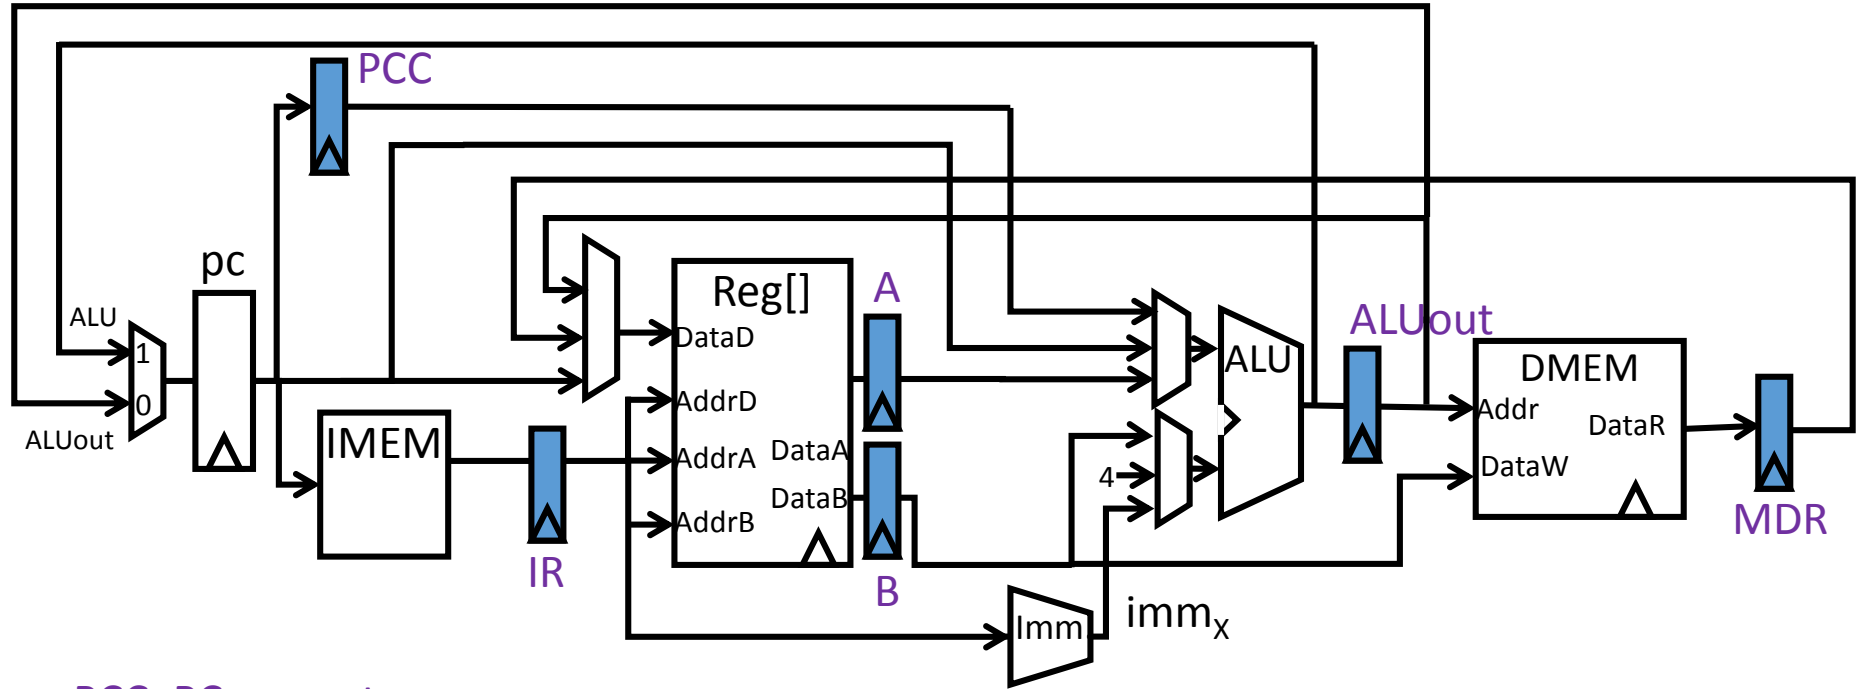
5 stages → 5 cycles (and loop back)



The controller is FSM



Adding temporary registers (and making some changes)



PCC=PC current

IR=Instruction Register

A
B

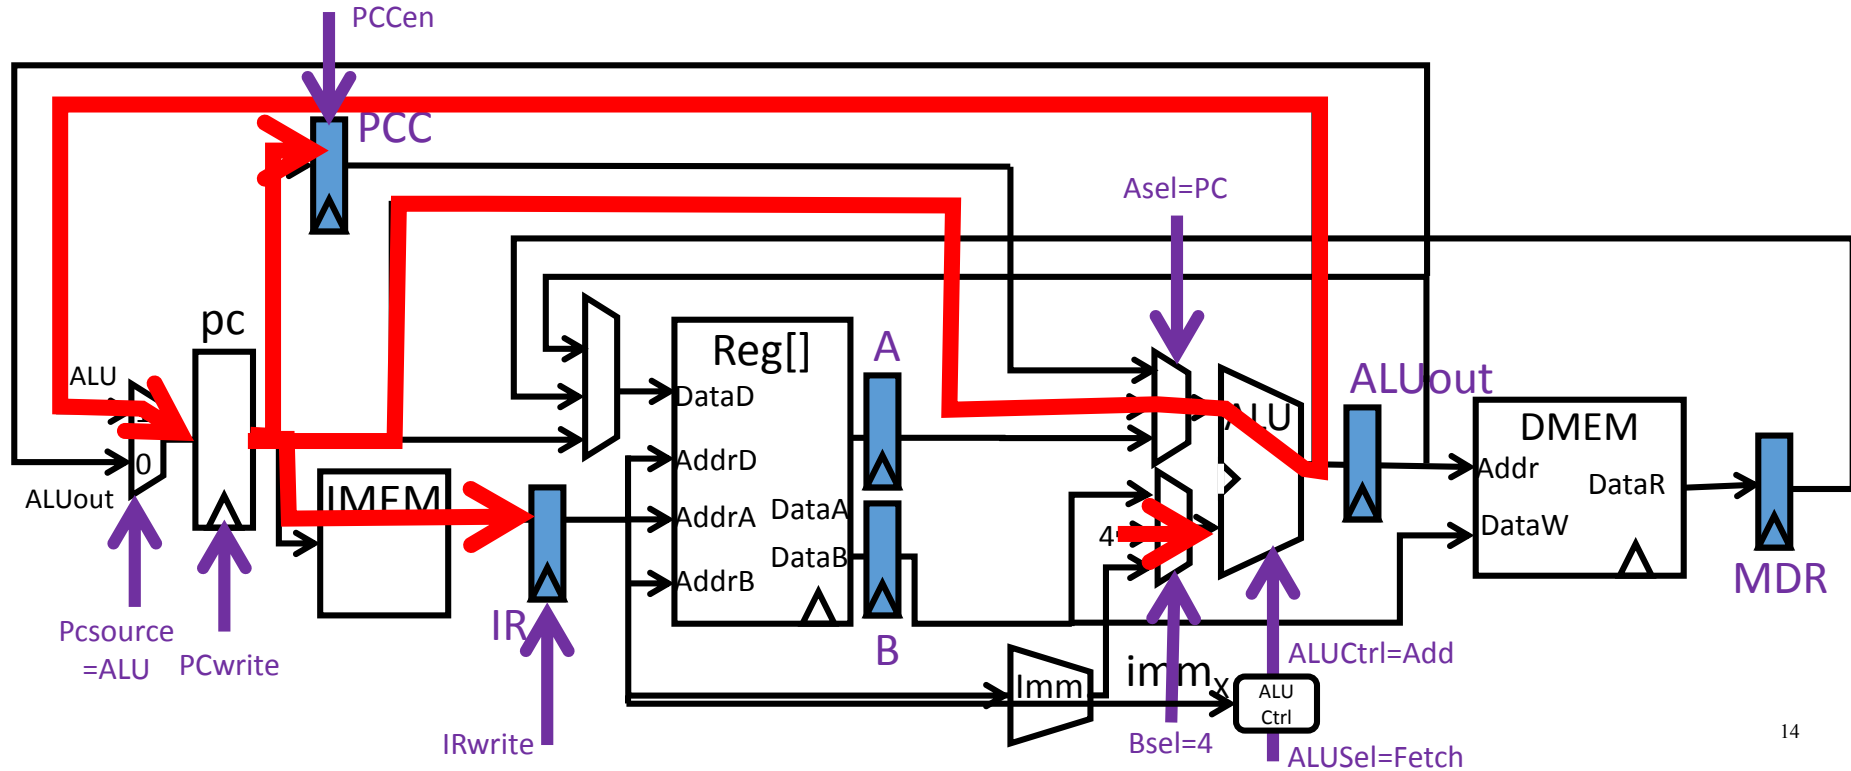
MDR=Memory Data Register

ALUout

FETCH cycle is the same for all instructions

FETCH: fetching the instruction pointed to by the PC

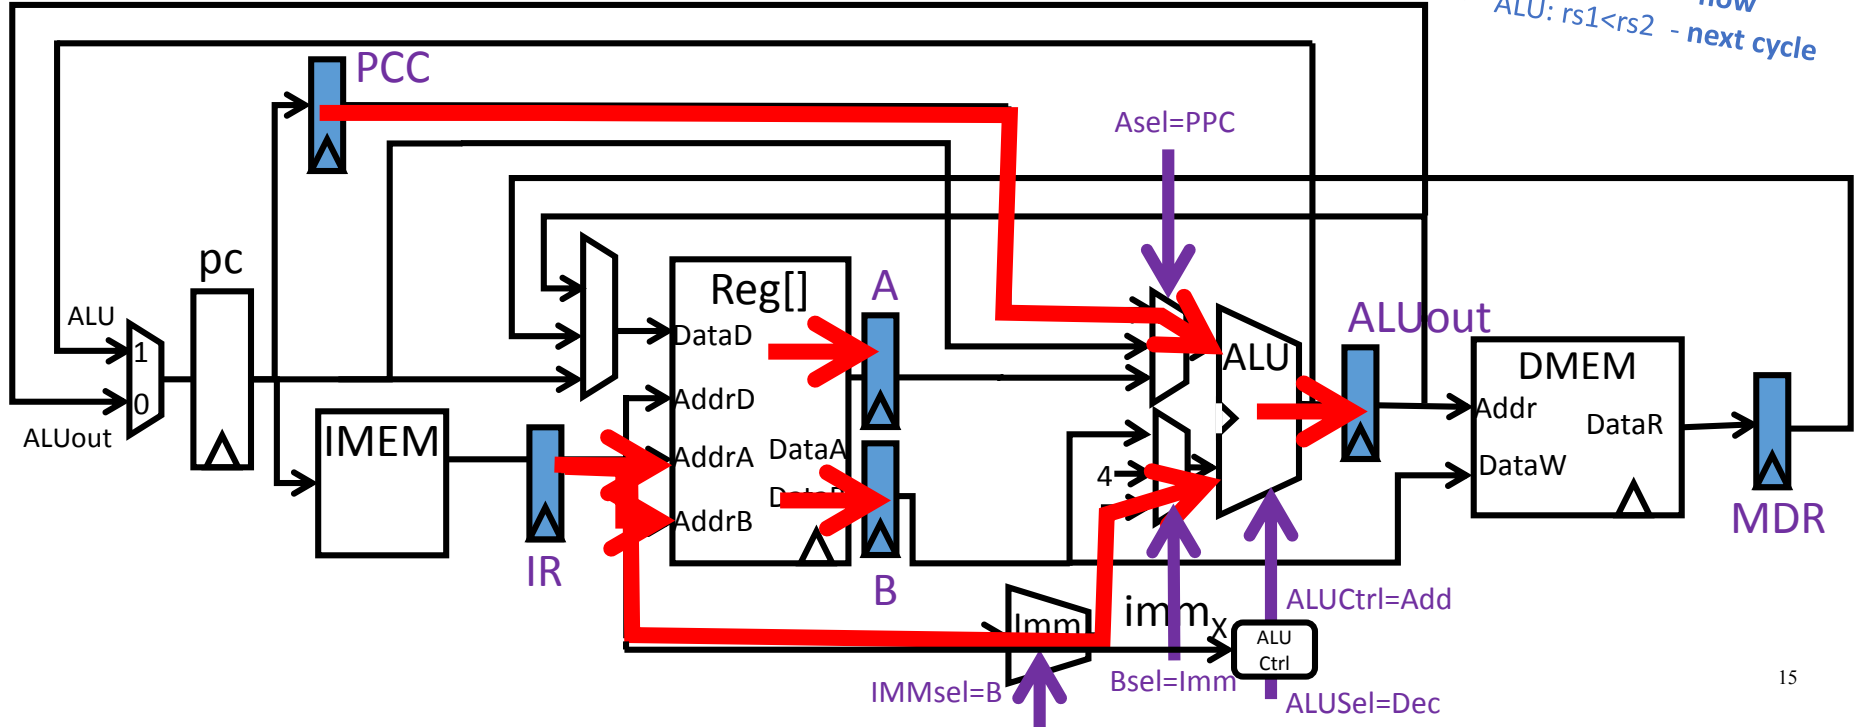
1. $IR = \text{Memory}[PC]$
2. $PC = PC + 4$, $PCC = PC$



DECODE cycle is the same for all instructions

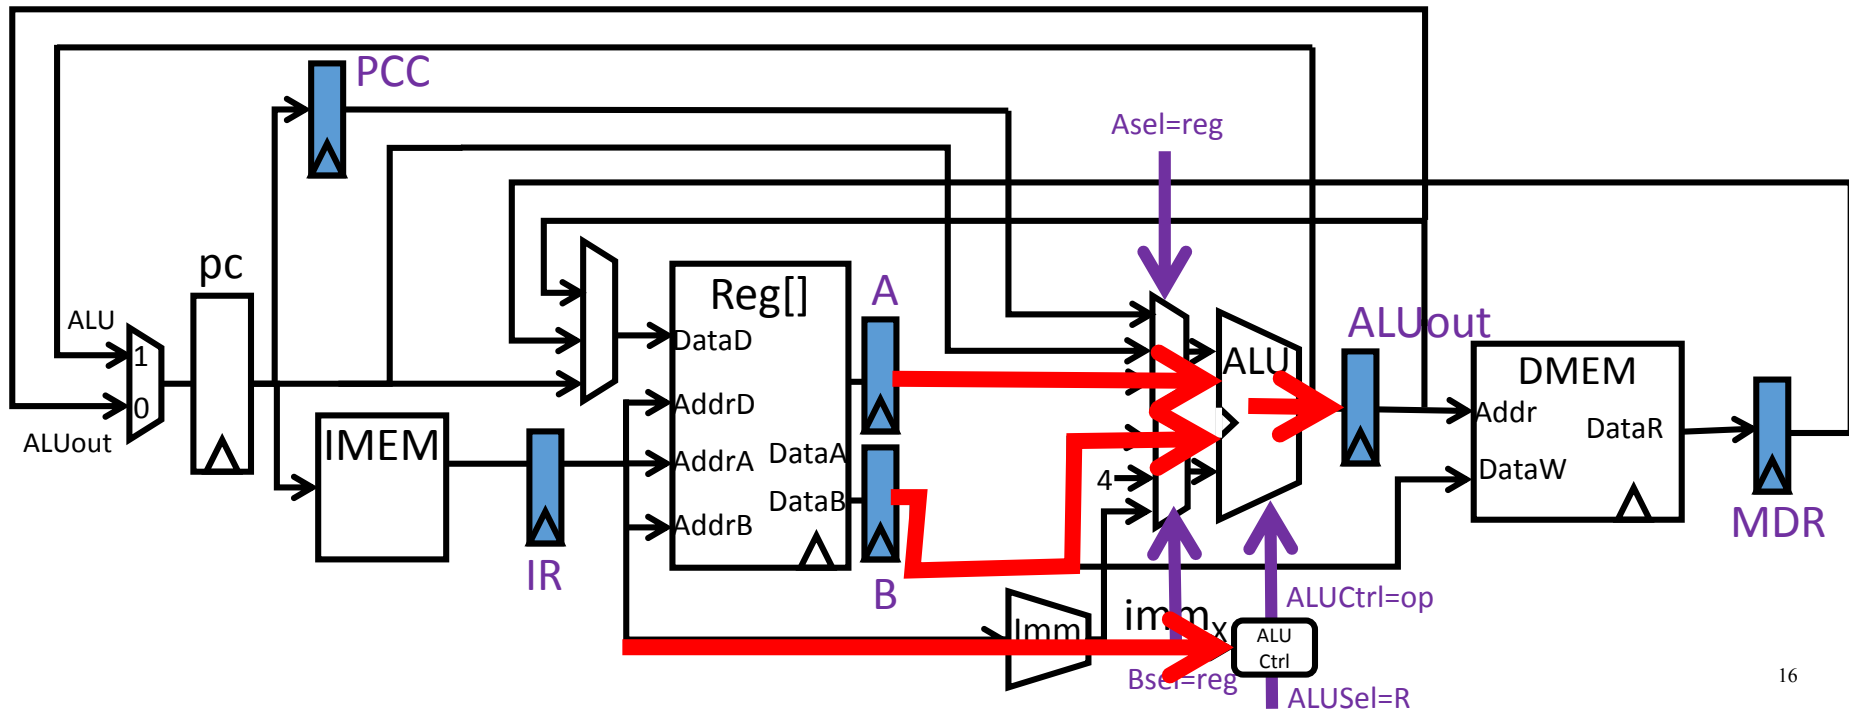
DECODE: decode instruction, read two registers, (speculatively) compute branch target address

1. $A = \text{Reg}[\text{rs1}], B = \text{Reg}[\text{rs2}]$
2. $\text{ALUOut} = \text{PCC} + \text{branch offset}$



Execute Cycle for R-Type

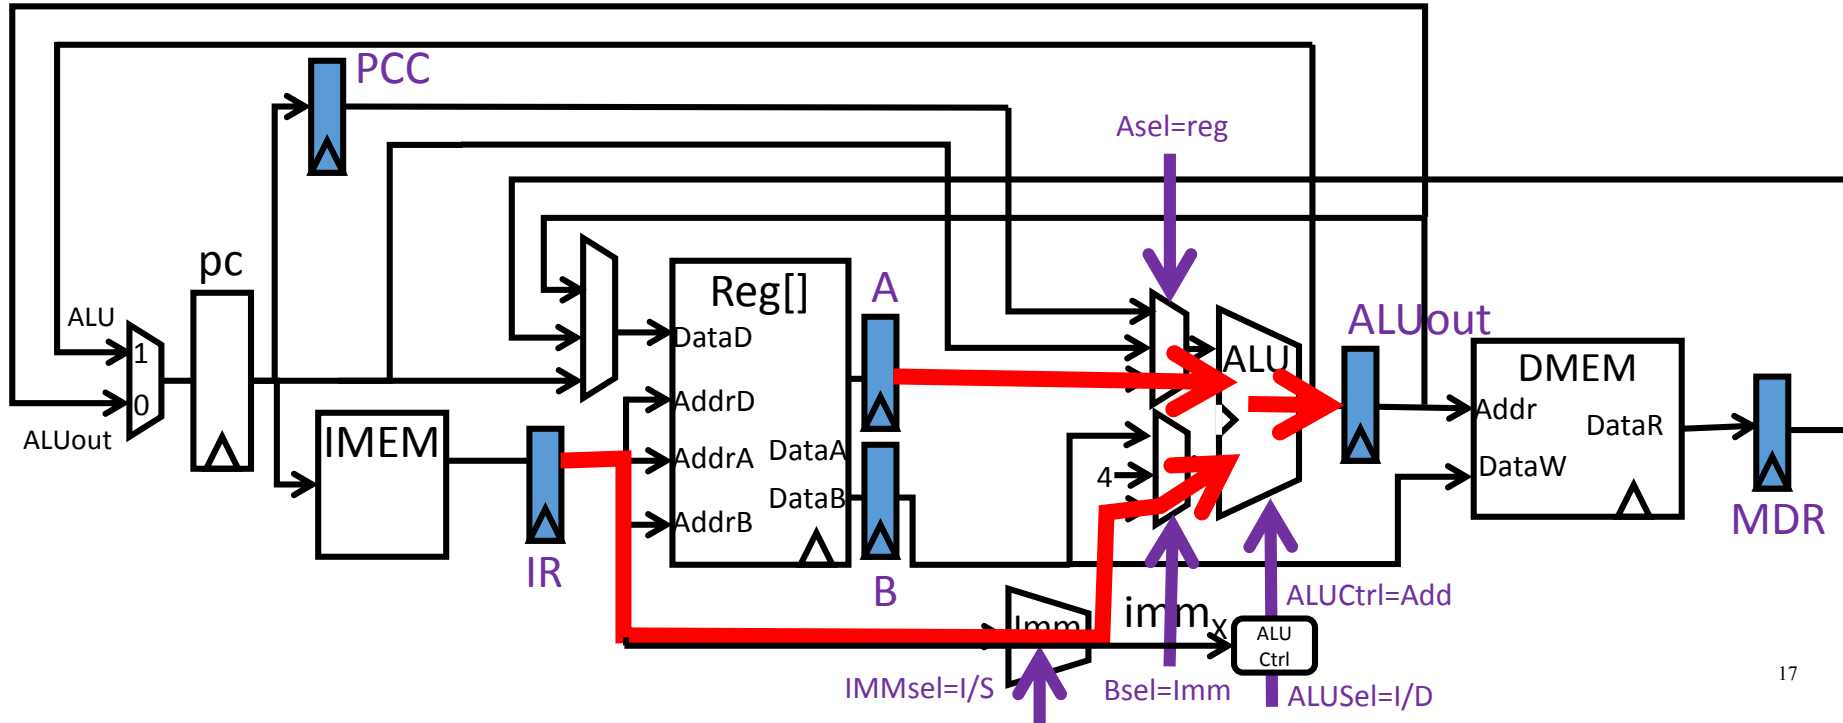
- $ALUOut = A \text{ op } B$



Execute Cycle for Load / Store

- ALUOut = A + immediate (I or S format)

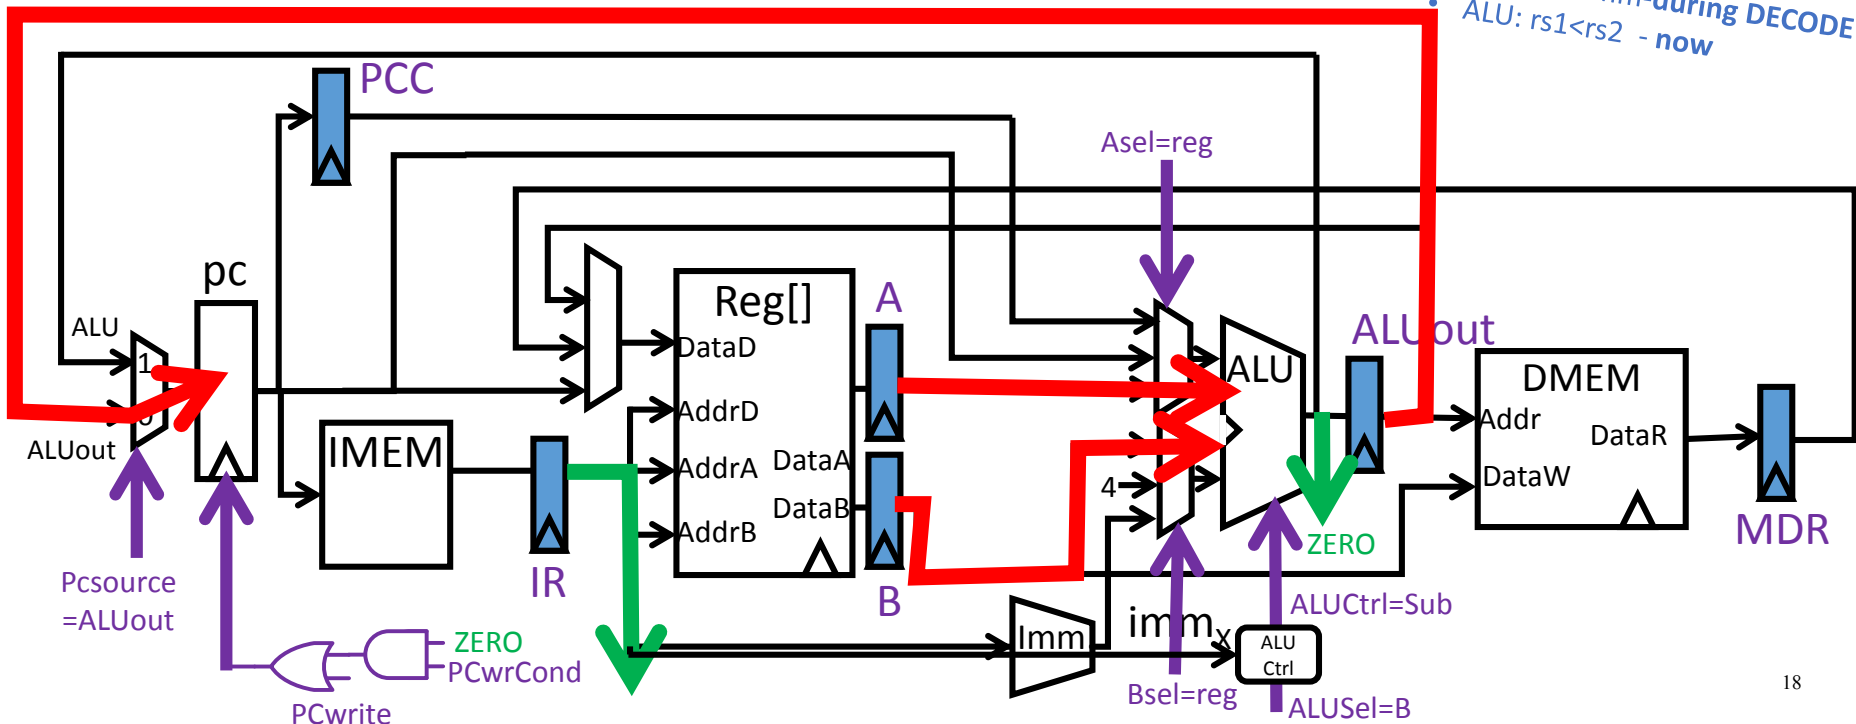
LW rd, rs1, imm
SW rs1, rs2, imm



Execute Cycle for Branch

- If (condition) then PC = ALUOut
- Branch condition is checked by ALU and controller (Zero output from ALU)

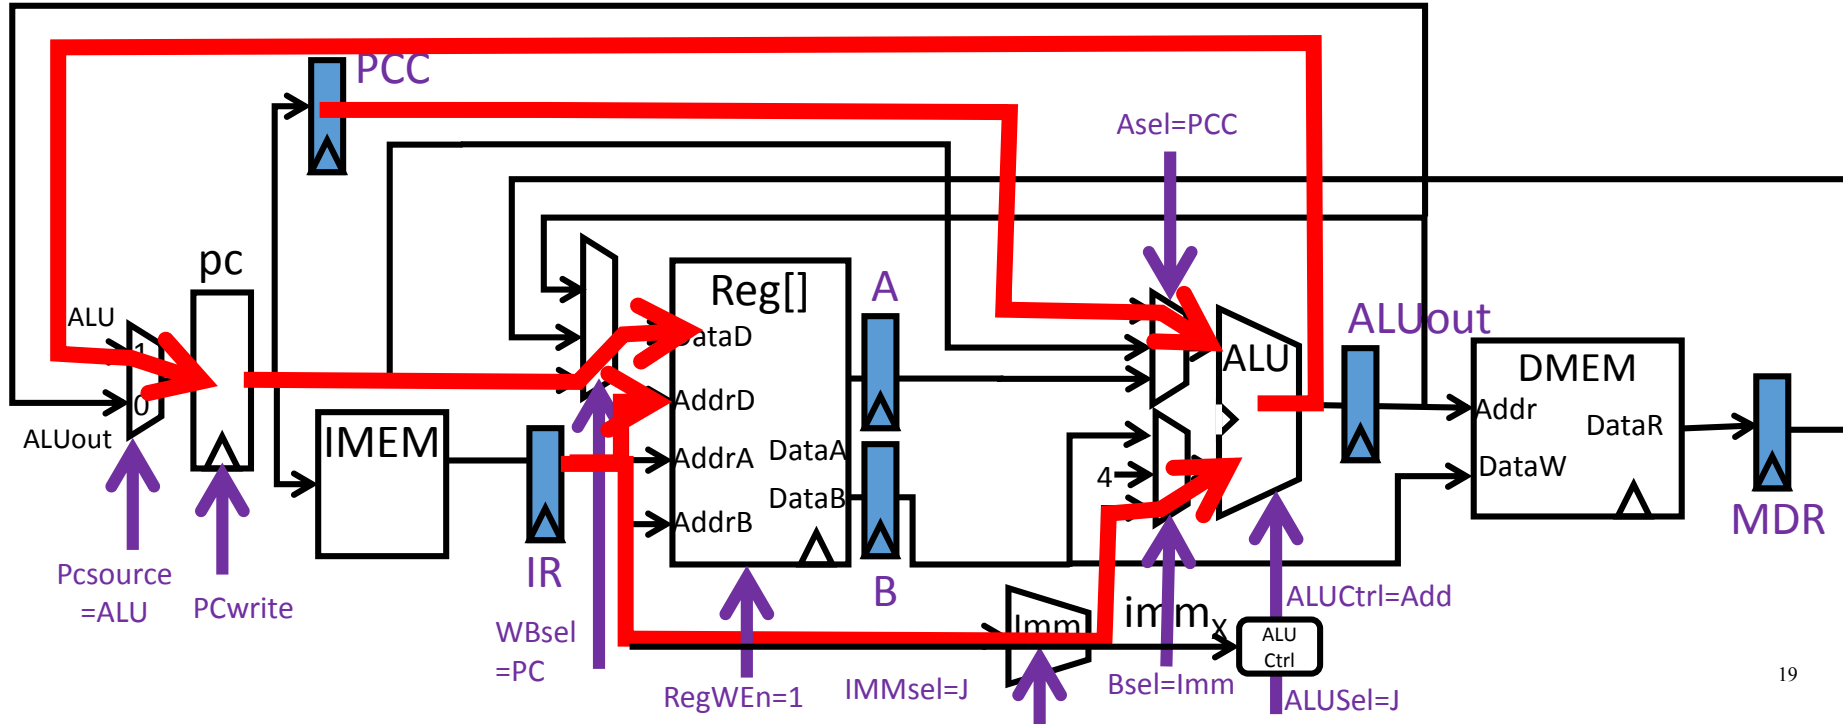
BLT rs1,rs2,imm
 • ALU: PC+imm-during DECODE
 • ALU: rs1<rs2 - now



Execute Cycle for JAL

- $\text{Reg}[\text{rd}] = \text{PC} + 4$ (the return address)
- $\text{PC} = \text{PCC} + \text{offset}$ (PC-relative jump)

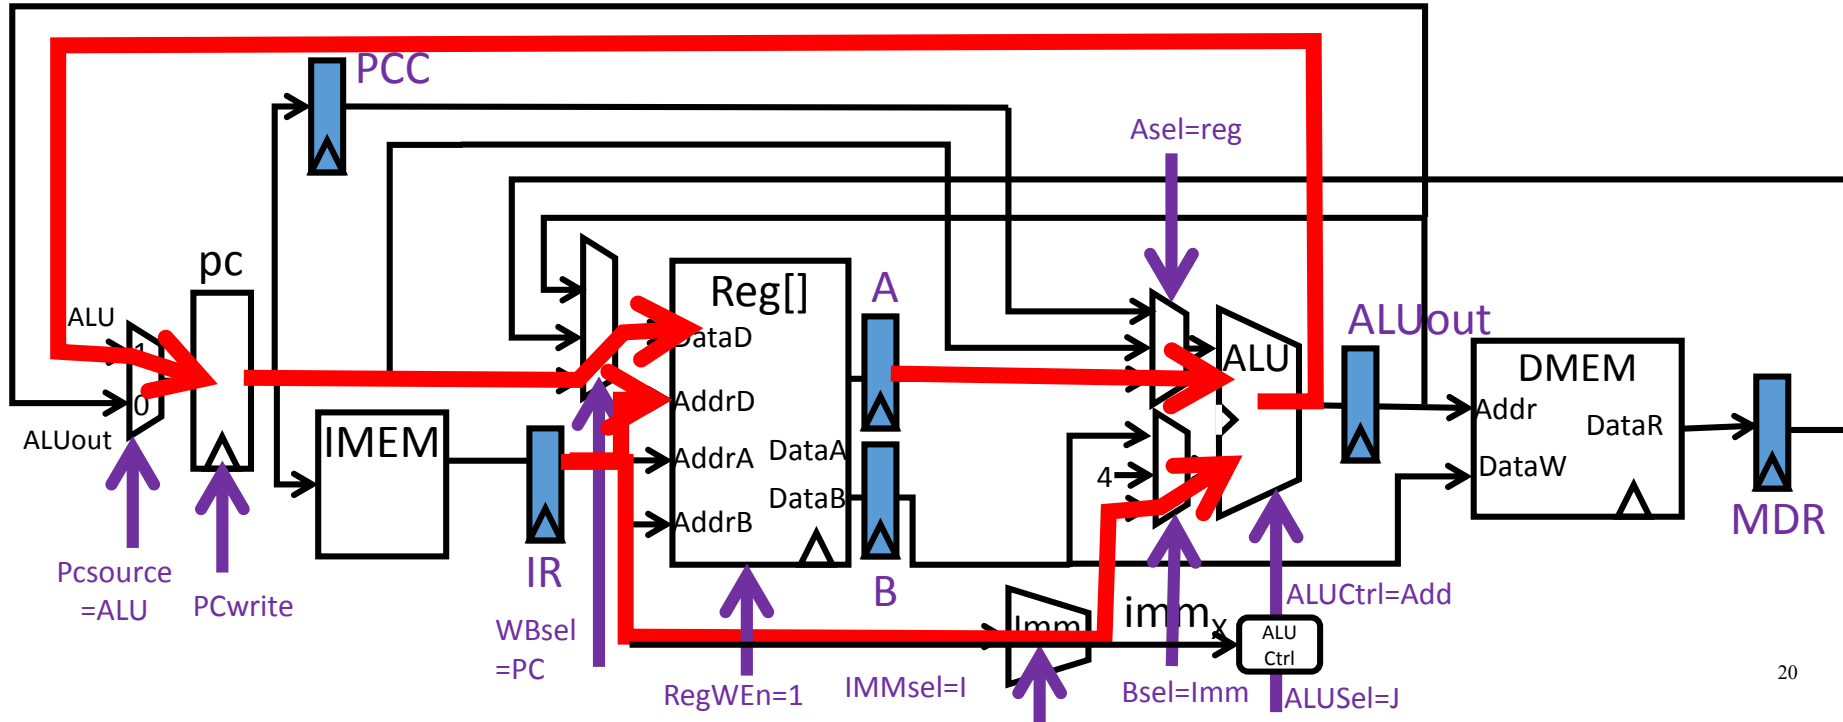
Recall: PC contains PC+4 after the Fetch cycle!



Execute Cycle for JALR

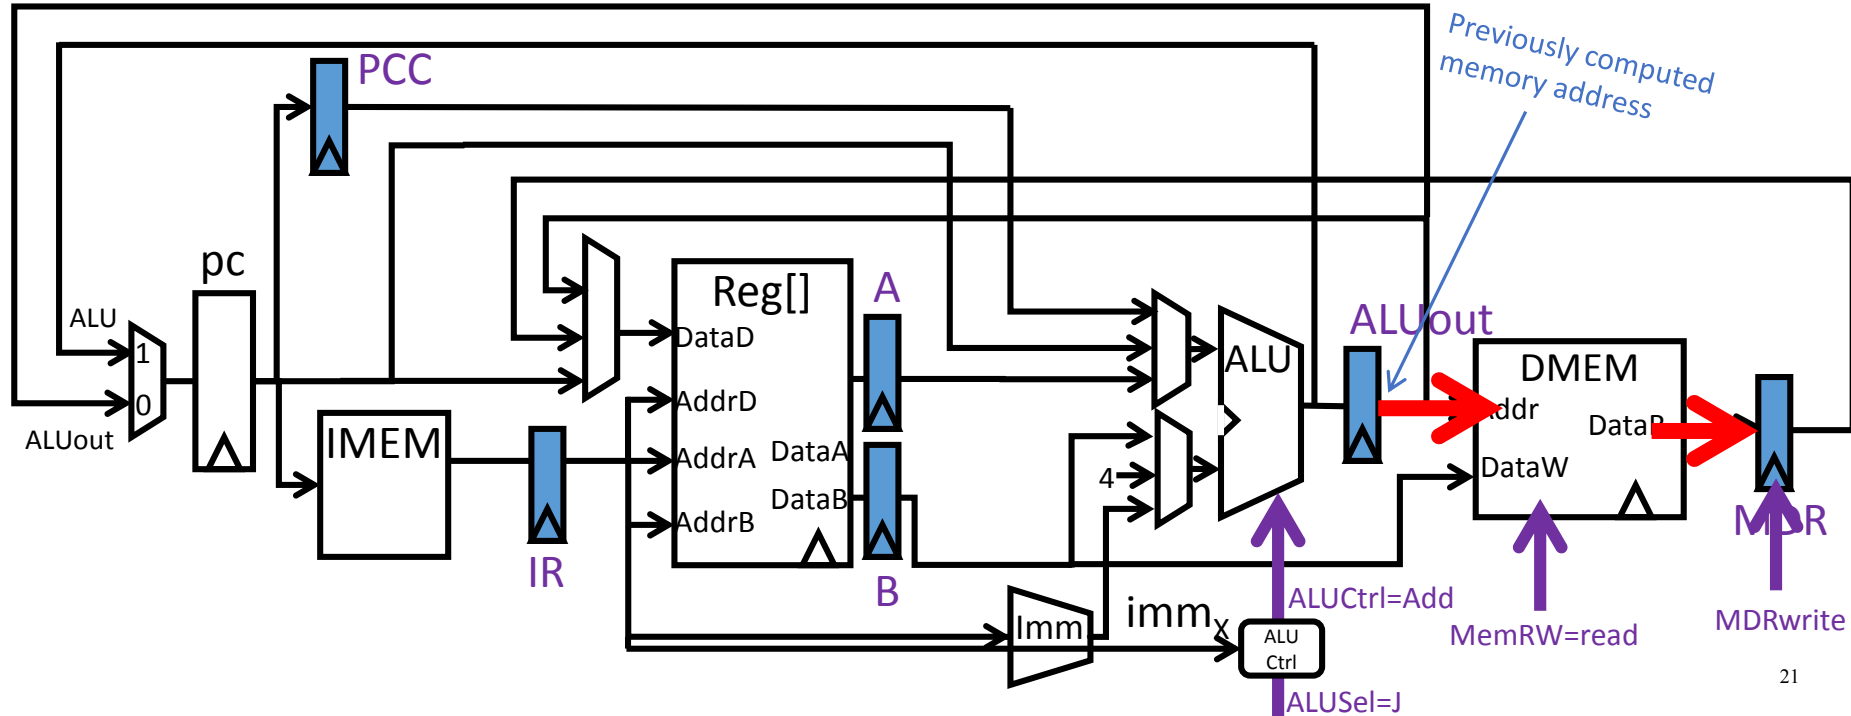
- $\text{Reg}[\text{rd}] = \text{PC} + 4$ (return address)
- $\text{PC} = \text{Reg}[\text{rs1}] + \text{immediate}$

JALR rd,rs1,imm



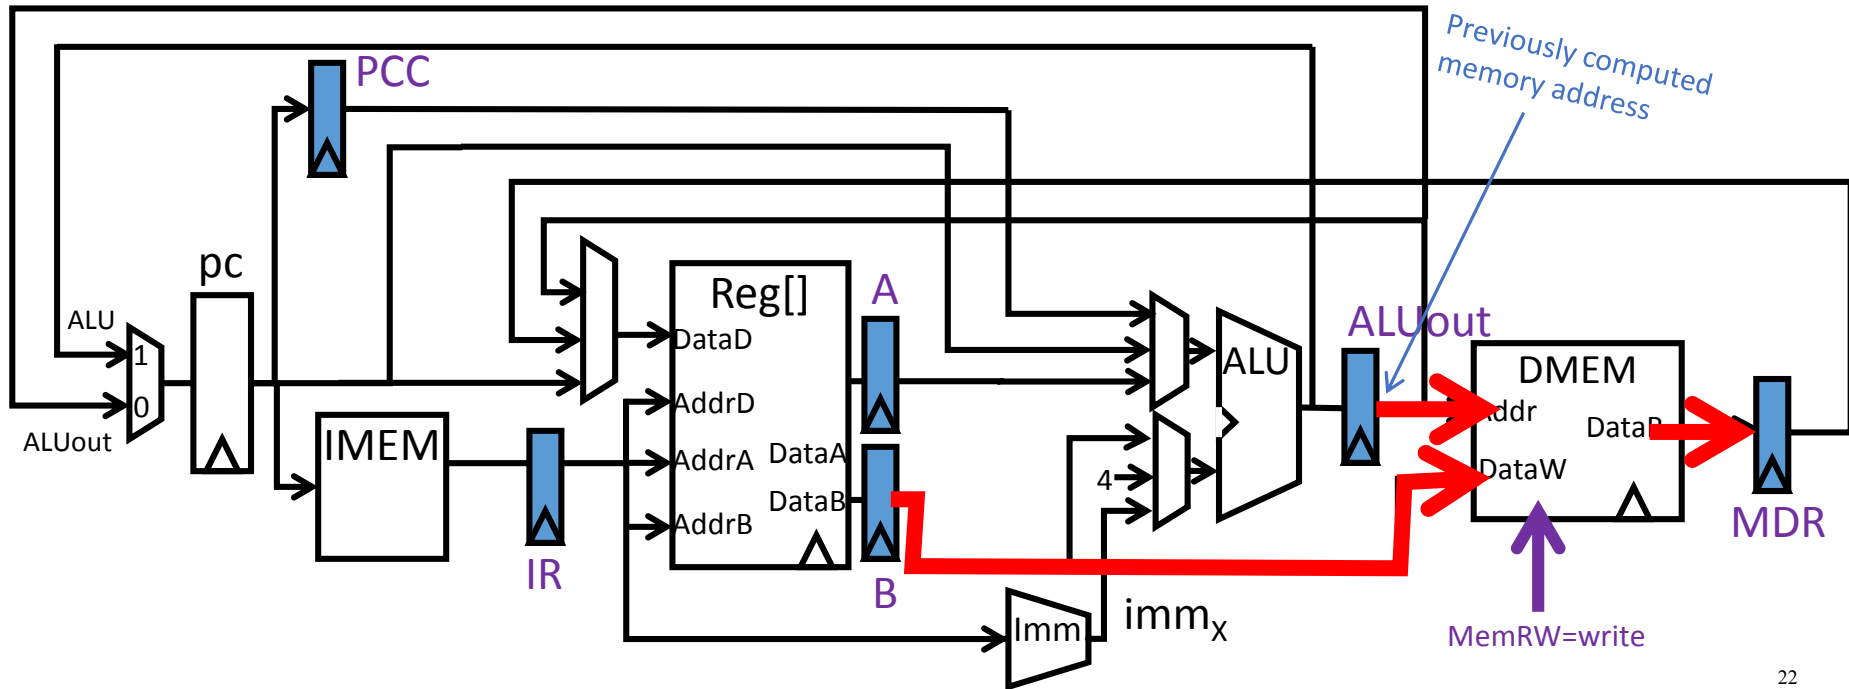
Memory cycle for Load

- MDR = Memory[ALUOut]



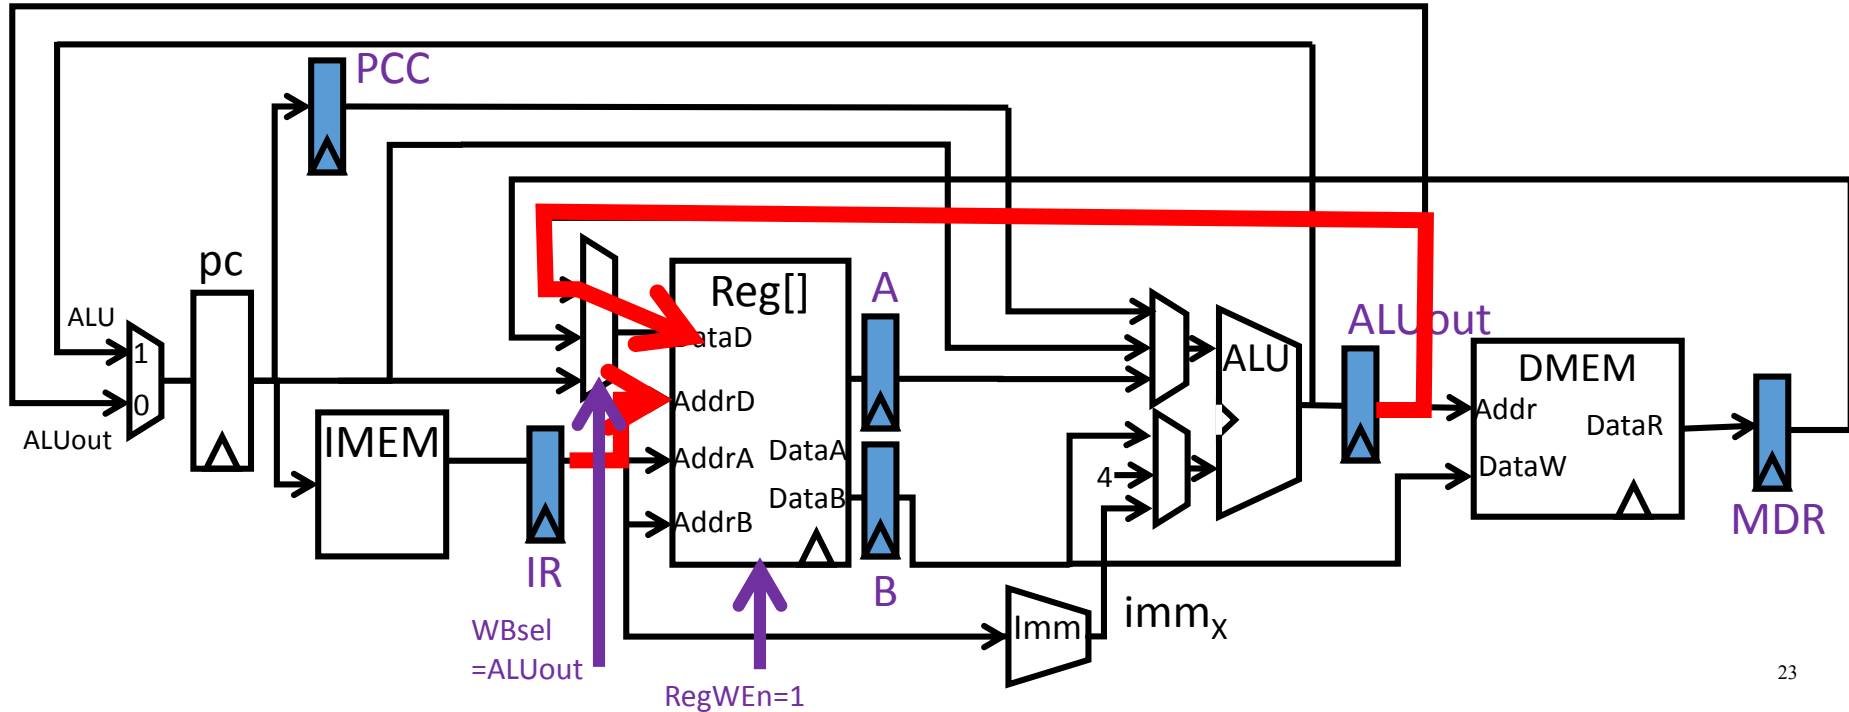
Memory cycle for Store

- $\text{Memory}[\text{ALUOut}] = B$



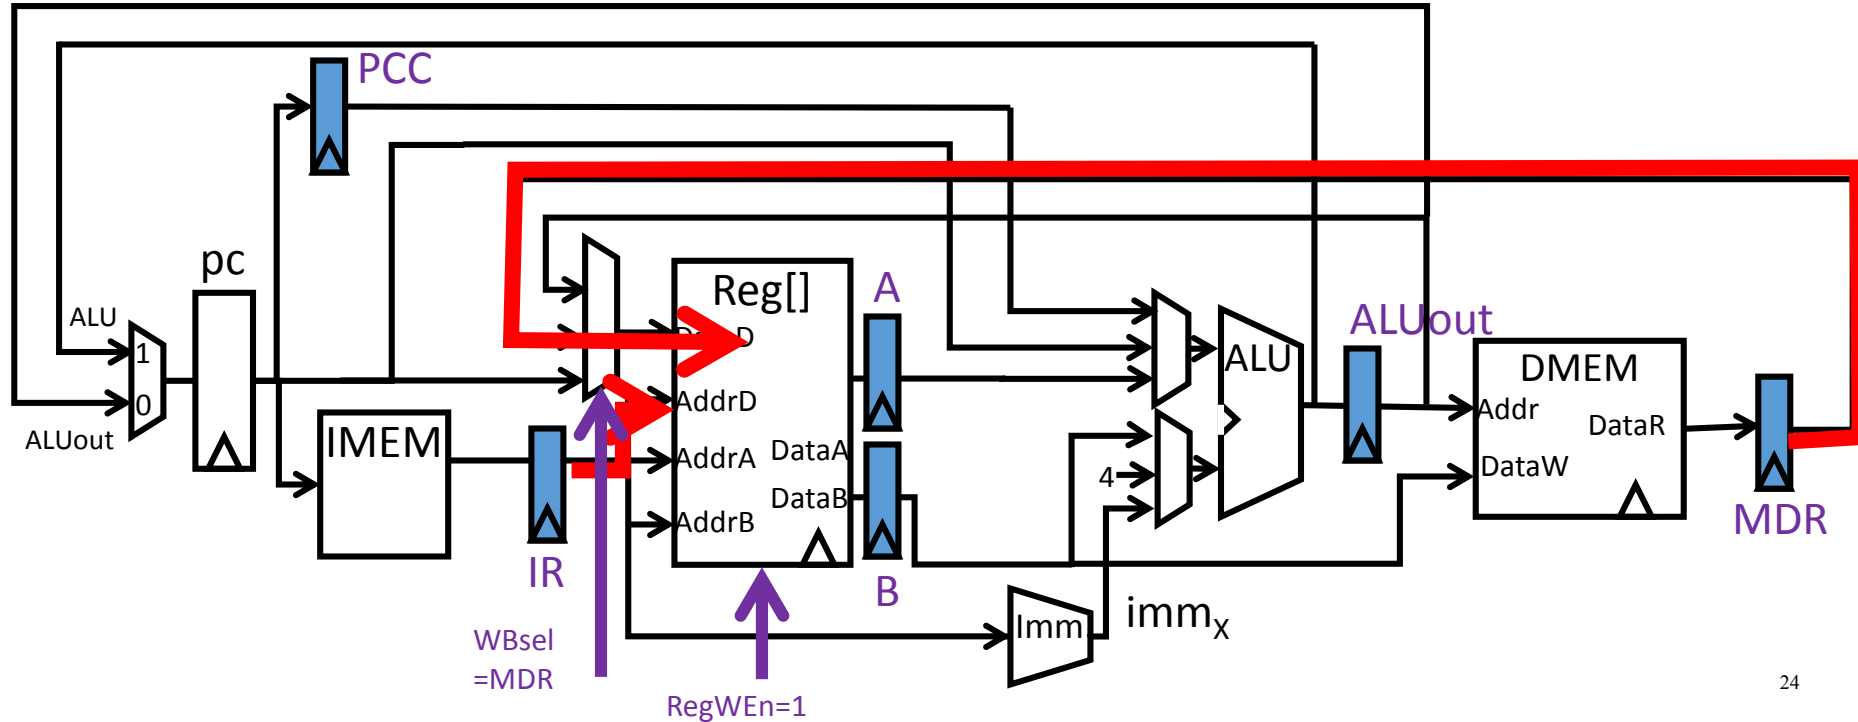
Write-back cycle for R-Type

- $\text{Reg}[\text{rd}] = \text{ALUOut}$



Write-back cycle for Load

- $\text{Reg}[\text{rd}] = \text{MDR}$



Summary

Step name	Action for R-type	Action for Load Instruction	Action for Store Instruction	Action for branches	Action for JAL
Instruction fetch	$IR = \text{Memory}[PC]$ $PC = PC + 4, PCC = PC$				
Instruction decode / operand fetch	$A = \text{Reg}[rs1]$ $B = \text{Reg}[rs2]$ $ALUOut = PCC + \text{branch offset}$				
Execution, address computation, branch/ jump completion	$ALUOut = A \text{ op } B$	$ALUOut = A + \text{immediate}$	$ALUOut = A + \text{immediate}$	if $(A == B)$ then $PC = ALUOut$	$\text{Reg}[rd] = PC$ $PC = PCC + \text{immediate}$
Memory access or R-type write-back	$\text{Reg}[rd] = ALUOut$	$MDR = \text{Memory}[ALUOut]$	$\text{Memory}[ALUOut] = B$		
Load write-back		$\text{Reg}[rd] = MDR$			

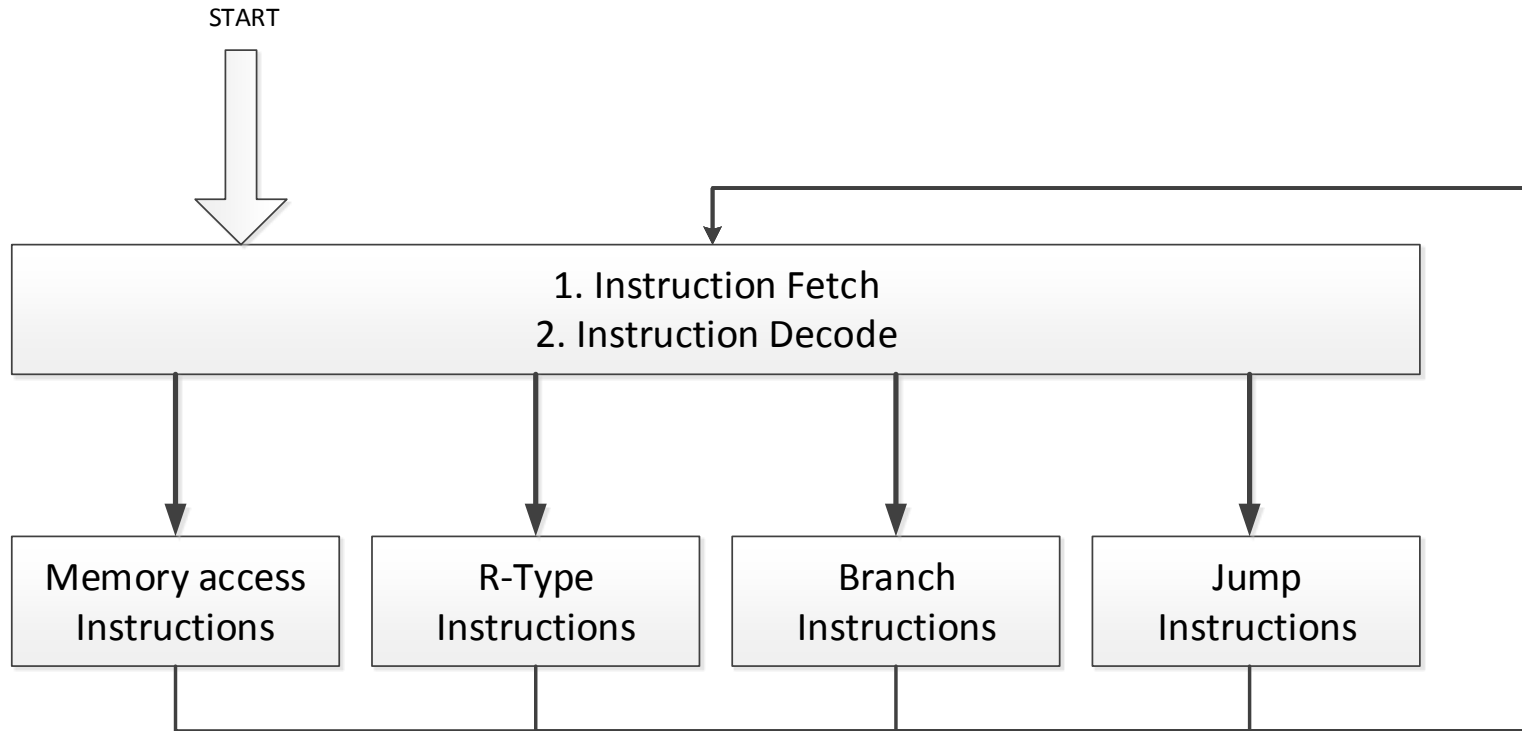
Example Questions

- How many cycles are needed to execute this code ?

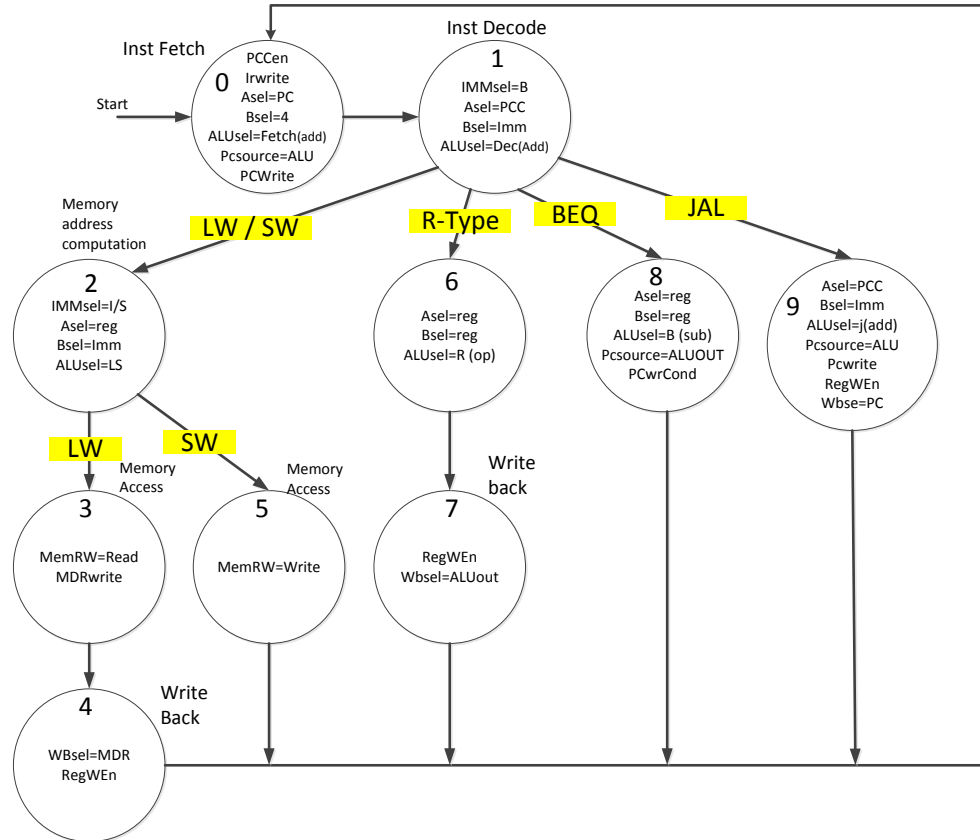
```
lw    $s2, 0($s3)
lw    $s3, 4($s3)
beq   $s2, $s3, Label #assume not
add   $s5, $s2, $s3
sw    $s5, 8($s3)
Label: ...
```

- What happens during the 8th cycle of execution ?
- In which cycle does the adding of $\$s2 + \$s3$ take place ?

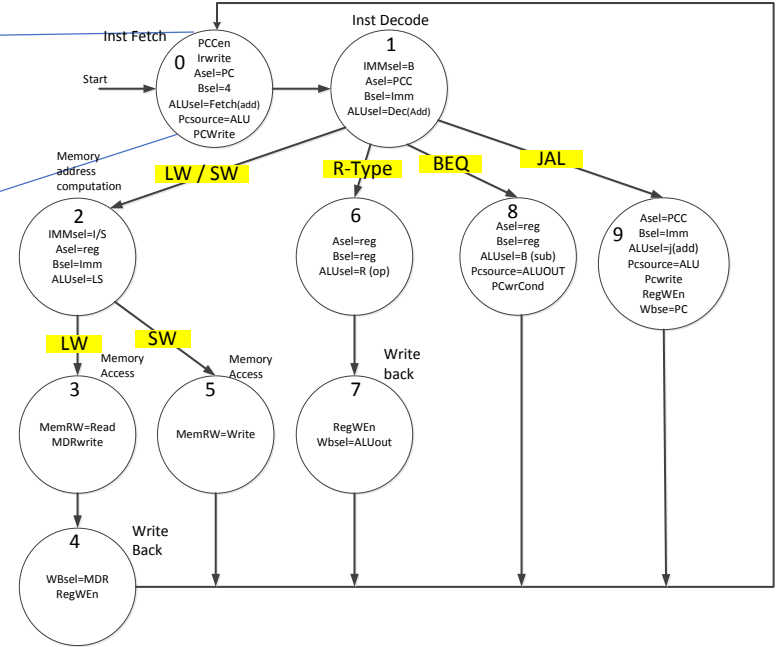
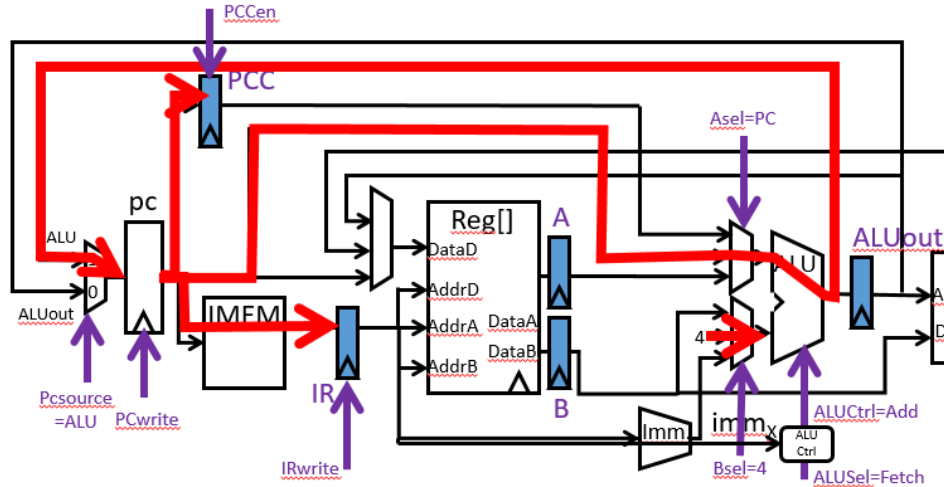

Controller State Machine



Controller State Machine




Fetch Cycle

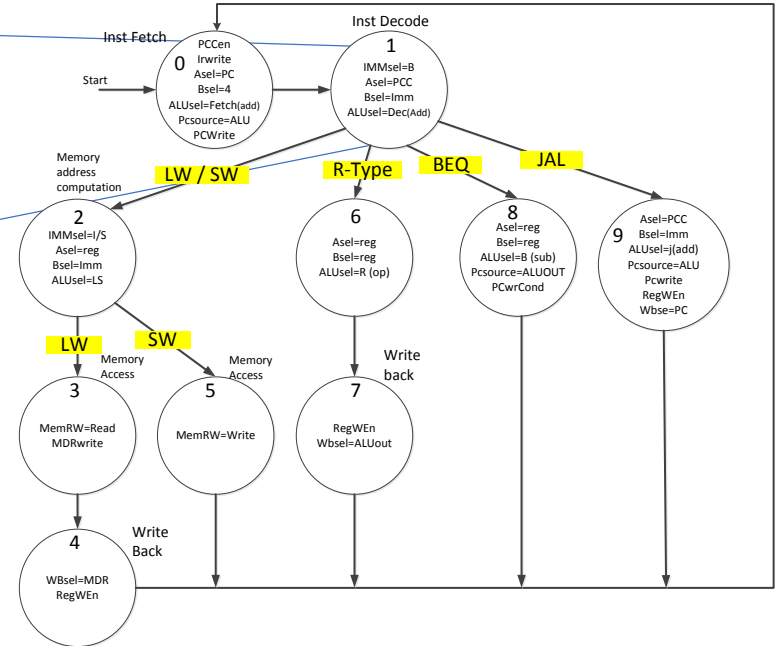
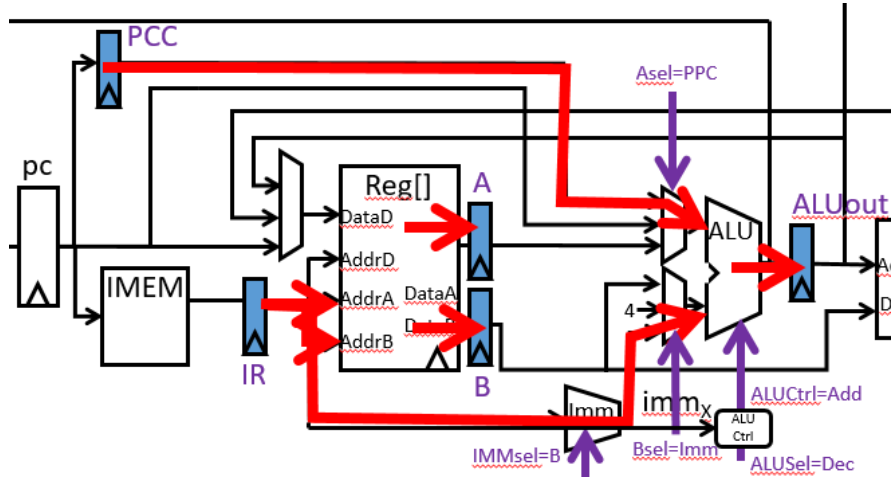


Instruction Decode / Register Fetch

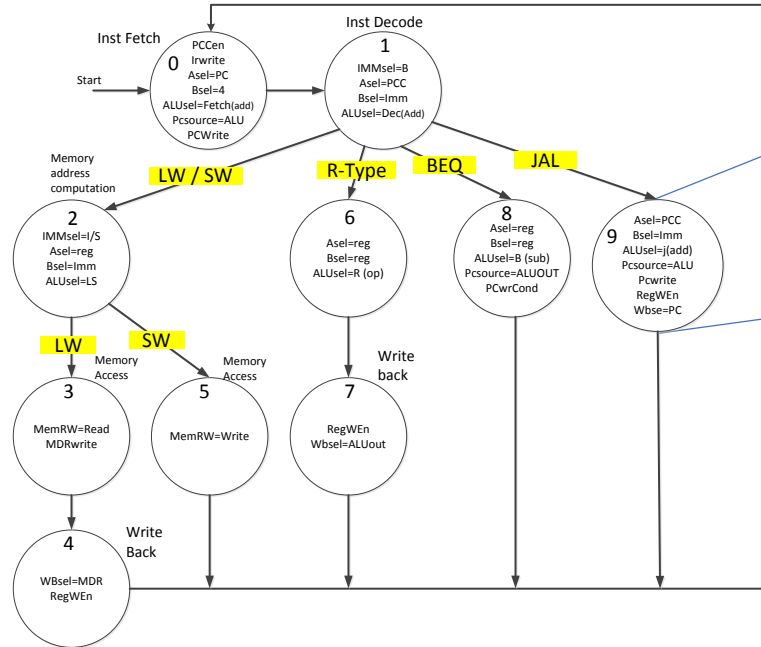
Inst Decode



1. $A = \text{Reg}[\text{rs1}], B = \text{Reg}[\text{rs2}]$
2. $\text{ALUOut} = \text{PCC} + \text{branch offset}$

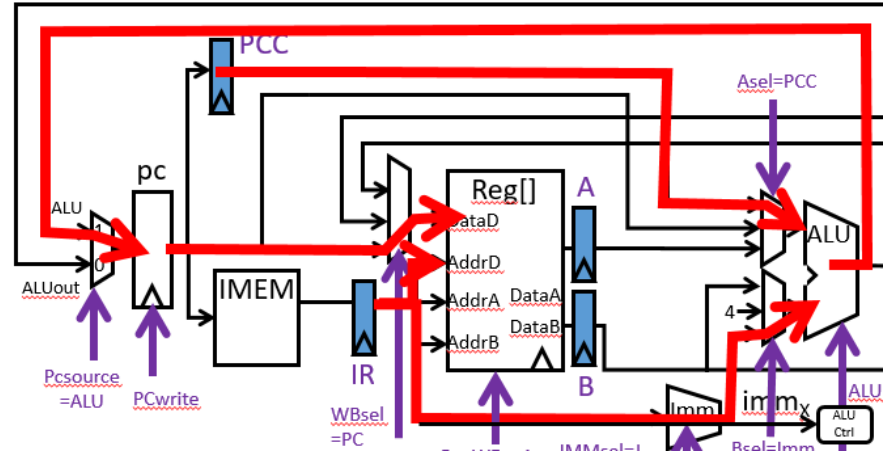


Execute cycle for JAL

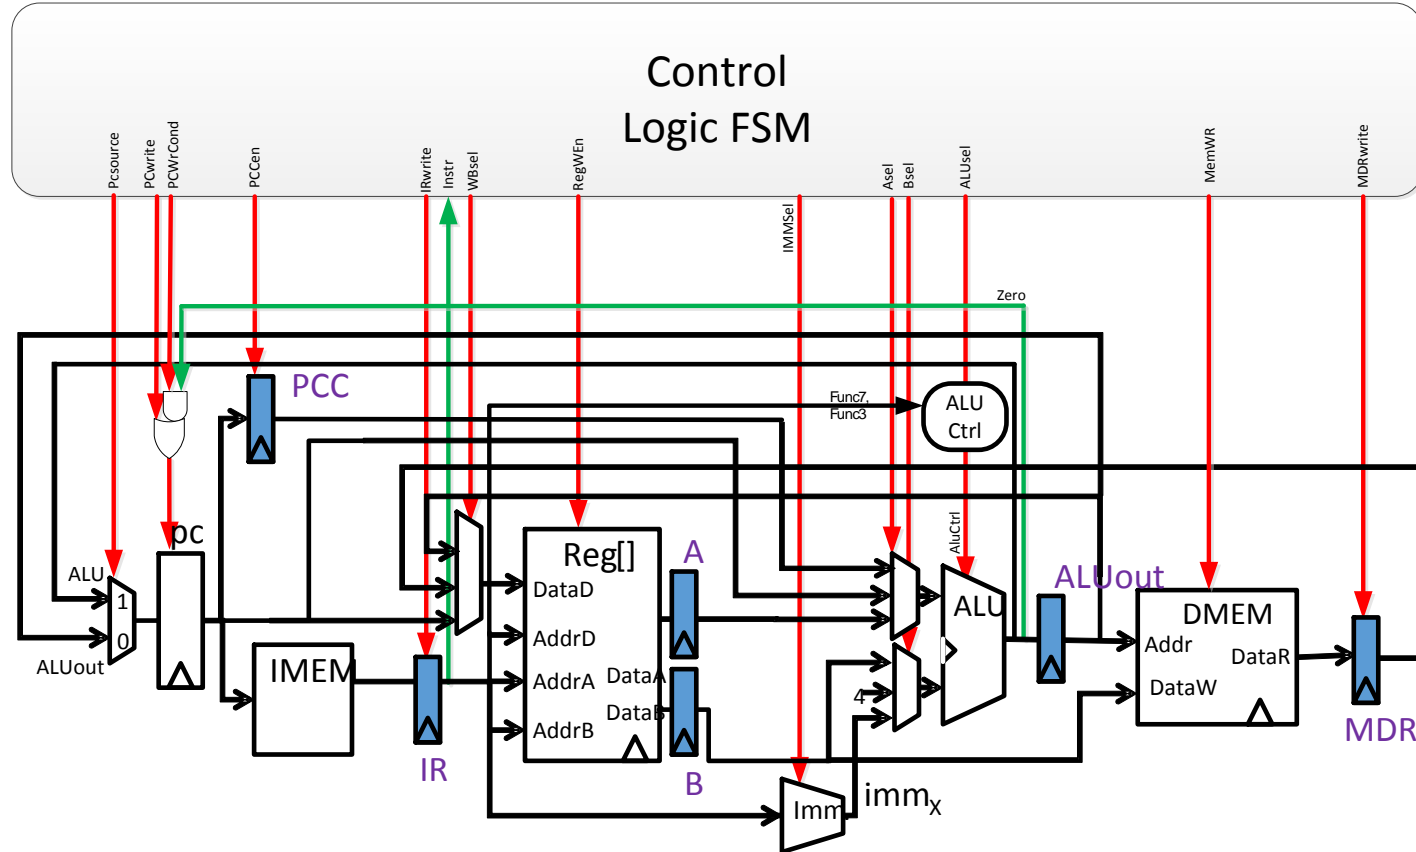


9
 Asel=PCC
 Bsel=Imm
 ALUset=j(add)
 Pcsource=ALU
 PCwrite
 RegWen
 WBsel=PC

- $\text{Reg}[\text{rd}] = \text{PC} + 4$ (the return address)
- $\text{PC} = \text{PCC} + \text{offset}$ (PC-relative jump)



Multi-Cycle RISC-V Controller and Datapath



Agenda

- Multi-Cycle RISC-V Datapath
- **Performance**
- ROM Controller
- Micro-Coded Controller

Performance evaluation: Cycles-Per-Instruction

- In the single-cycle RISC-V, $CPI=1$
- In the multi-cycle RISC-V, CPI depends on instruction:
 - Load: 5 cycles
 - Store: 4 cycles
 - ALU: 4 cycles
 - Branch: 3 cycles
 - Jump: 3 cycles


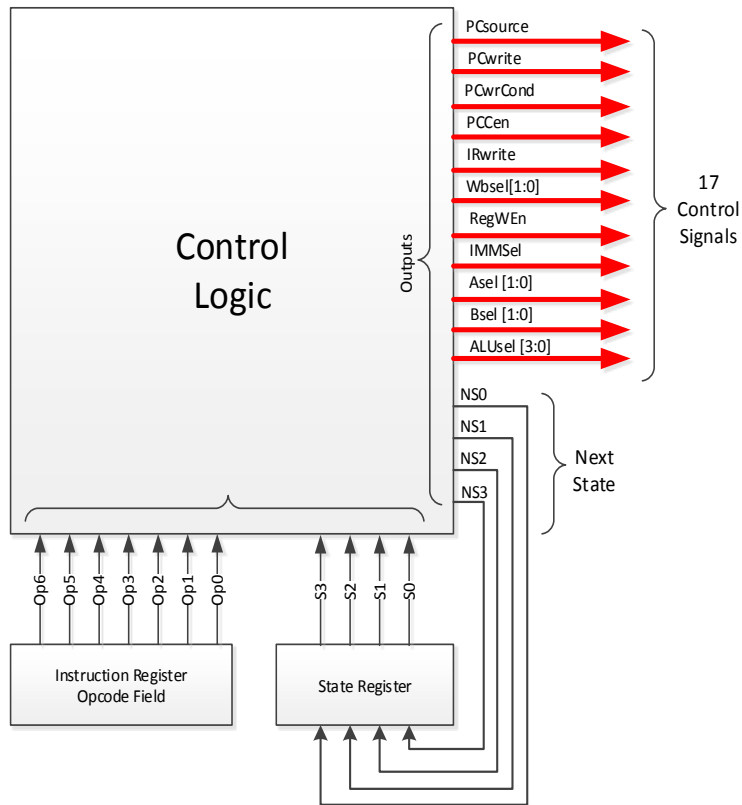
Performance improvement due to multi-cycle

- By instruction dynamic distribution (see slide on Elastic Clock Cycle)
 - $\text{CPI} = 5 \times 24\% + 4 \times 12\% + 4 \times 44\% + 3 \times 18\% + 3 \times 2\% = 4.02$ cycles/instruction (on average)
- 4.02 is better than worst case of 5 (all instructions use 5 cycles)
- BUT apparently single-cycle RISC-V is “better” – $\text{CPI}=1$...
- Not really. Instruction TIME matters: $\text{Instruction Time} = \text{Cycle time} \times \text{CPI}$
- What was Cycle Time in single-cycle RISC-V ? 800 ps
- What is Cycle Time in multi-cycle RISC-V ?
 - We split execution to five parts. Ideally, $\text{Cycle Time} = 800 \text{ ps} / 5 = 160 \text{ ps}$
 - But practice is never ideal. Hard to split exactly equal
 - Extra registers may also add certain overhead (setup time, clock-to-output time)
- If Cycle Time is indeed $800/5=160$ ps, then $\text{Instruction Time} = \text{Cycle time} \times \text{CPI} = 160\text{ps} \times 4.02 = 643.2 \text{ ps}$
- In this ideal case, improvement is $800/643.2 = 5/4.02 = \underline{1.24}$ (multi-cycle 24% faster than single cycle)

Agenda

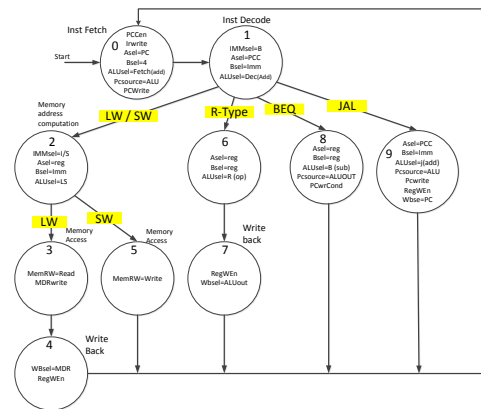
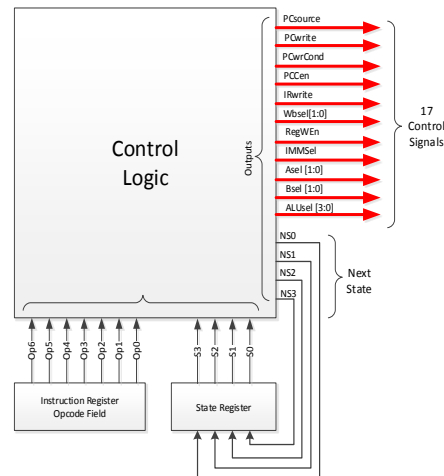
- Multi-Cycle RISC-V Datapath
- Performance
- **ROM Controller**
- Micro-Coded Controller

Implementing the FSM



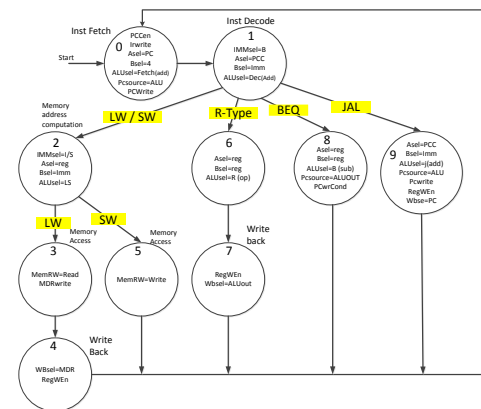
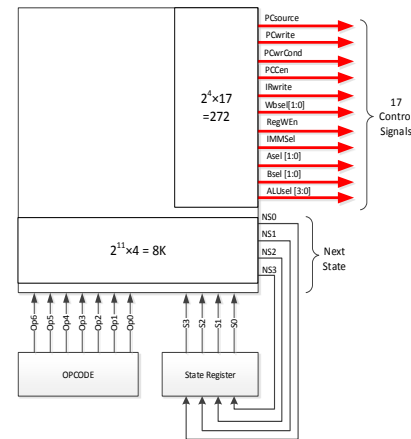
Controller implemented as a ROM

- Inputs
 - 7 opcode bits
 - 4 state bits
 - Total 11 bits $\rightarrow 2^{11}=2048$ words
- Outputs
 - 17 control bits
 - 4 state bits
 - Total 21 bits per word
- ROM size $2048 \times 21 = \underline{42 \text{ Kbits}}$




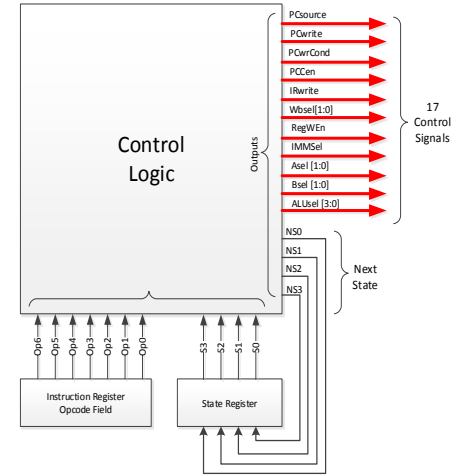
Reducing ROM size

- Break ROM into two
 - Moore FSM \rightarrow 17 control outputs are function of only 4 state bits:
 $(2^4 \text{ words}) \times (17 \text{ bits/word}) = 272 \text{ bits}$
 - 11 inputs (7 op + 4 state) determine next state:
 $(2^{11} \text{ words}) \times (4 \text{ bits/word}) = 8\text{K bits}$
 - Total 8.25 Kbits (cf. 42 Kbits)


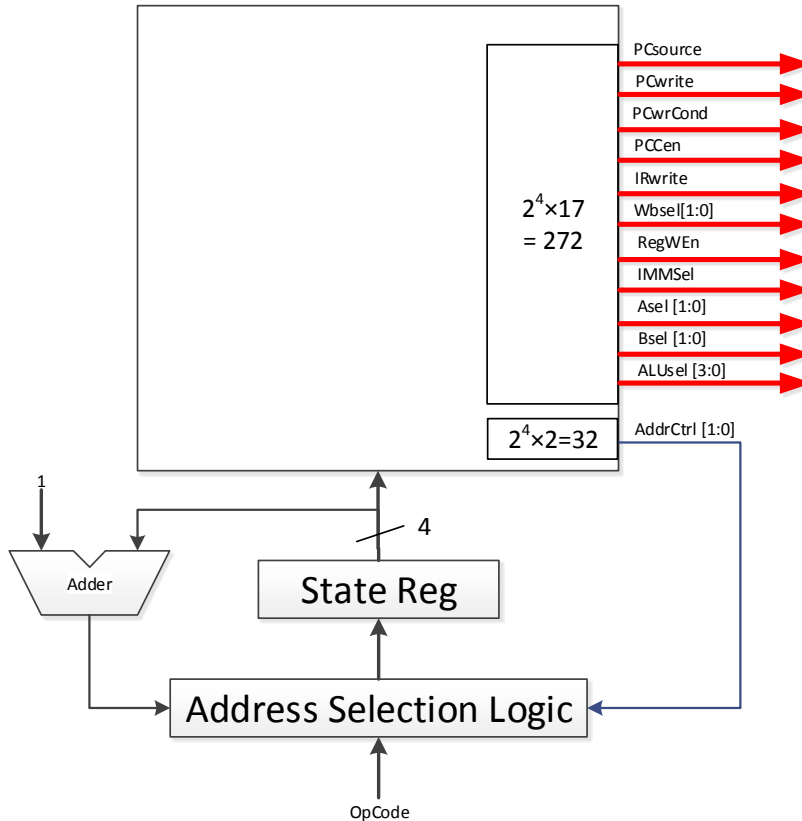


Optimizing FSM

- ROM is non-optimal
 - Contains all combinations, some of them useless
 - Often we know exactly what next state is
 - Often, next state is next word in the ROM
 - In these cases, we don't need the next state bits
- We could have minimized the logic and use gates....
- Better idea: Use counter, increment +1 for next state (often)
- Often, but not always. So add a new control signal (use counter or not) [**AddrCtrl**]
- VERY similar to the PC
 - Often go to PC+4, otherwise jump/branch




Added counter for next state


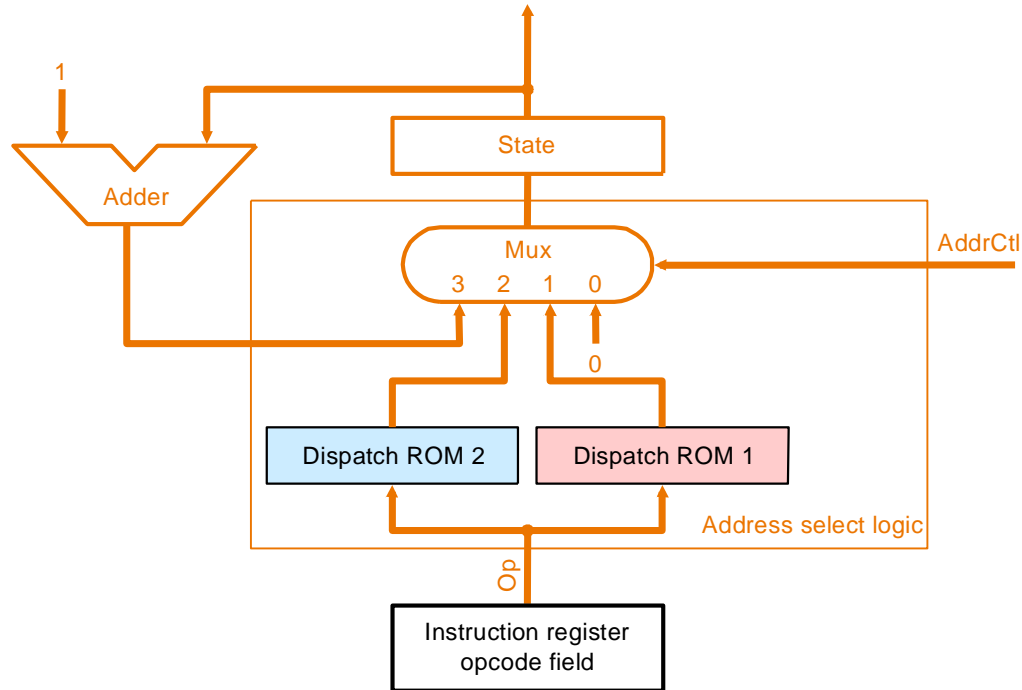


What if next state is not state+1 ?

- State diagram branches in 1, 2
- Use branch table for 1
- Use branch table for 2
- Go to zero after 4,5,7,8,9
- Go to state+1 after 0,3,6




Address Selection



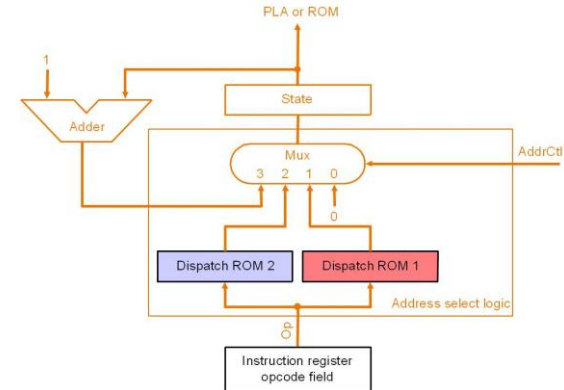
Address Selection Tables

- טבלה 1 לקפיצה ממצב מספר 1
- טבלה 2 לקפיצה ממצב מספר 2



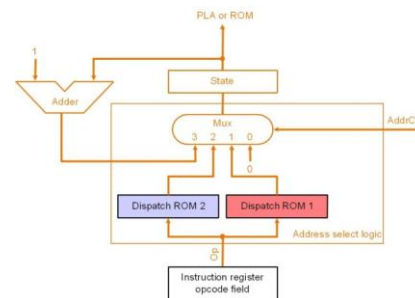

Dispatch ROM 1		
Op	Name	Value
x33	R-type	6
x67	JAL	9
x63	BEQ	8
x03	LW	2
x23	SW	2

Dispatch ROM 2		
Op	Name	Value
x03	LW	3
x23	SW	5



Address Selection Control

State number	Address control action	Value of AddrCtl
0	Increment	3
1	Use dispatch ROM 1	1
2	Use dispatch ROM 2	2
3	Increment	3
4	Go to state 0	0
5	Go to state 0	0
6	Increment	3
7	Go to state 0	0
8	Go to state 0	0
9	Go to state 0	0

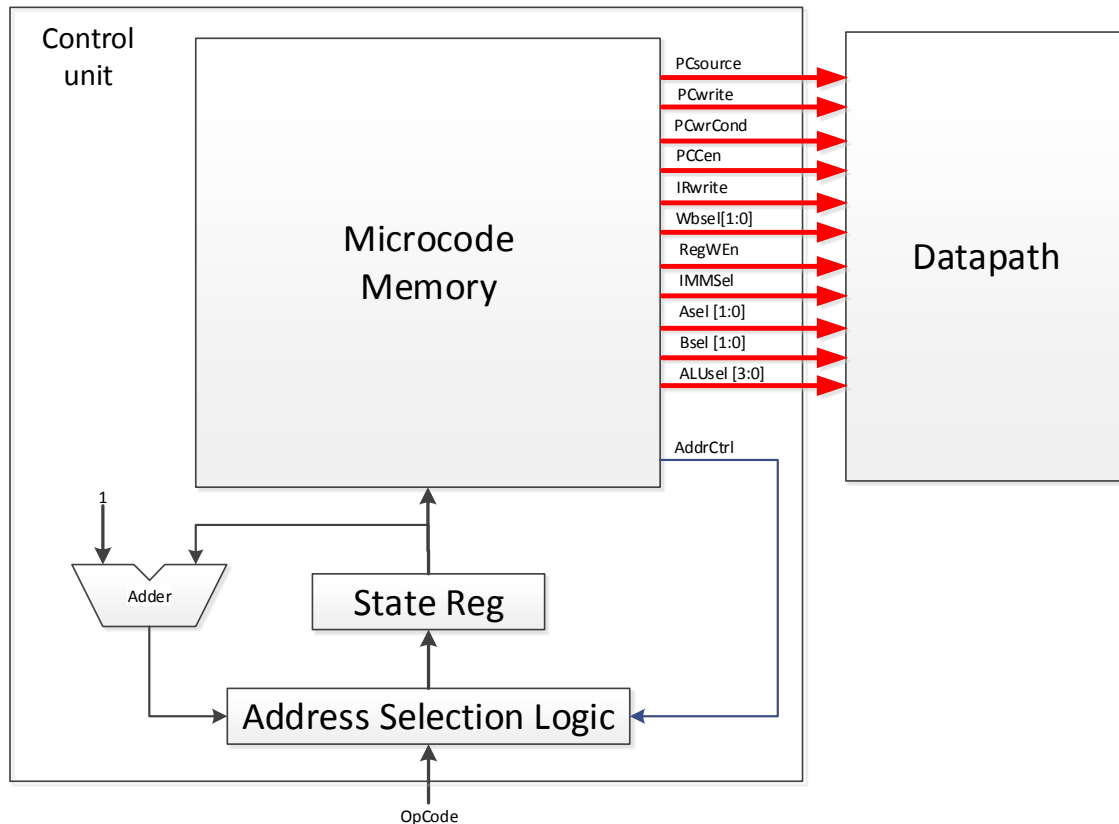


Result is smaller than ROM—no 4 bits/word for next state, only two bits

Agenda

- Multi-Cycle RISC-V Datapath
- Performance
- ROM Controller
- **Micro-Coded Controller**

Micro-code Controller



This is the modified ROM controller

Let's change ROM to RAM

Now we can change control RAM

This is Microcode / Microprogram

Micro-code


- μ PC is analogous to PC
- μ Code is analogous to Code
- μ Instruction is analogous to Instruction

- Multiple μ Instructions are (micro-) executed to implement one instruction
- Often, easier to make complex instructions by μ Code section than by state diagram
 - Large state diagrams are hard to read (hence, error prone)
 - μ Code is symbolic, translated by μ Assembler. No need to manage bits
- Most complex computers implemented that way

Micro-coded RISC-V controller

- Symbolic μ code
- Tabulated

State	Label	ALUsel	Asel	Bsel	MemRW	PCsource	WBsel	Sequencing
0	Fetch	Fetch	PC	4		ALU		Seq
1		Dec	PCC	Imm				Dispatch 1
2	Mem1	LS	Reg	Imm				Dispatch 2
3	LW2				Read			Seq
4							MDR	Fetch
5	SW2				Write			Fetch
6	Rformat1	R	Reg	Reg				Seq
7								Fetch
8	BEQ1	B	Reg	Reg		ALUout		Fetch
9	JAL	J	PCC	Imm		ALU	PC	Fetch



Note
This table is only partial

μInstruction Format (partial)

	Fetch	ALUCtrl = 000	Cause the ALU to add for Fetch stage.
	Dec	ALUCtrl = 001	Cause the ALU to add for decode stage.
ALU control	LS	ALUCtrl = 010	Cause the ALU to add for Load/Store instructions
	B	ALUCtrl = 011	Cause the ALU to subtract; this implements the compare for branches.
	J	ALUCtrl = 100	Cause the ALU to add for jump target address computation
	R	ALUCtrl = 101	Use the instruction's function code (Func3/7) to determine ALU control.
	...		
Asel	PC	Asel = 01	Use the PC as the first ALU input.
	PCC	Asel = 10	Use the PCC as the first ALU input.
	Reg	Asel = 00	Register A is the first ALU input.
	B	Bsel = 00	Register B is the second ALU input.
Bsel	4	Bsel = 01	Use 4 as the second ALU input.
	Imm	Bsel = 10	Use output of the sign extension unit as the second ALU input.
	Read	RegWEn=0	Read two registers using the rs1 and rs2 fields of the IR as the register numbers and putting the data into registers A and B.
	Write ALUout	RegWEn= 1, WBsel = 00	Write a register using the rd field of the IR as the register number and the contents of the ALUOut as the data.
Register control	Write MDR	RegWEn= 1, WBsel = 01	Write a register using the rd field of the IR as the register number and the contents of the MDR as the data.
	Write PC	RegWEn= 1, WBsel = 10	Write a register using the rd field of the IR as the register number and the contents of the PC as the data.
	Read DMEM	MemRW=0	Read memory using address calculated in EXE stage and sittinh in ALUout
		MDRwrite=1	Save the read data into MDR
Memory (DMEM)	Write DMEM	MemRW=1	Write to memory using the ALUOut as address and B register as data
		MDRwrite=0	
	ALU	PCSource = 1, PCWrite=1	Write the output of the ALU into the PC. For PC+4, JAL address, etc.
PC write control	ALUOut	PCSource = 0, PCWrite=1	If the Zero output of the ALU is active, write the PC with the contents of the register ALUOut.
	Seq	AddrCtl = 11	Choose the next microinstruction sequentially.
Sequencing	Fetch	AddrCtl = 00	Go to the first microinstruction to begin a new instruction.
	Dispatch 1	AddrCtl = 01	Dispatch using the ROM 1.

Conclusions

- Single cycle RISC-V
 - Lots of hardware (dedicated to functionality)
 - Slow
 - Combinational control
- Multi-cycle RISC-V
 - Hardware savings (but added registers)
 - Faster
 - Some speculation – more energy consumed
 - Sequential control (ROM / micro-code)
 - Flexible: see homework. We later exploit this to show interrupts